

# **Memory Management Basics**

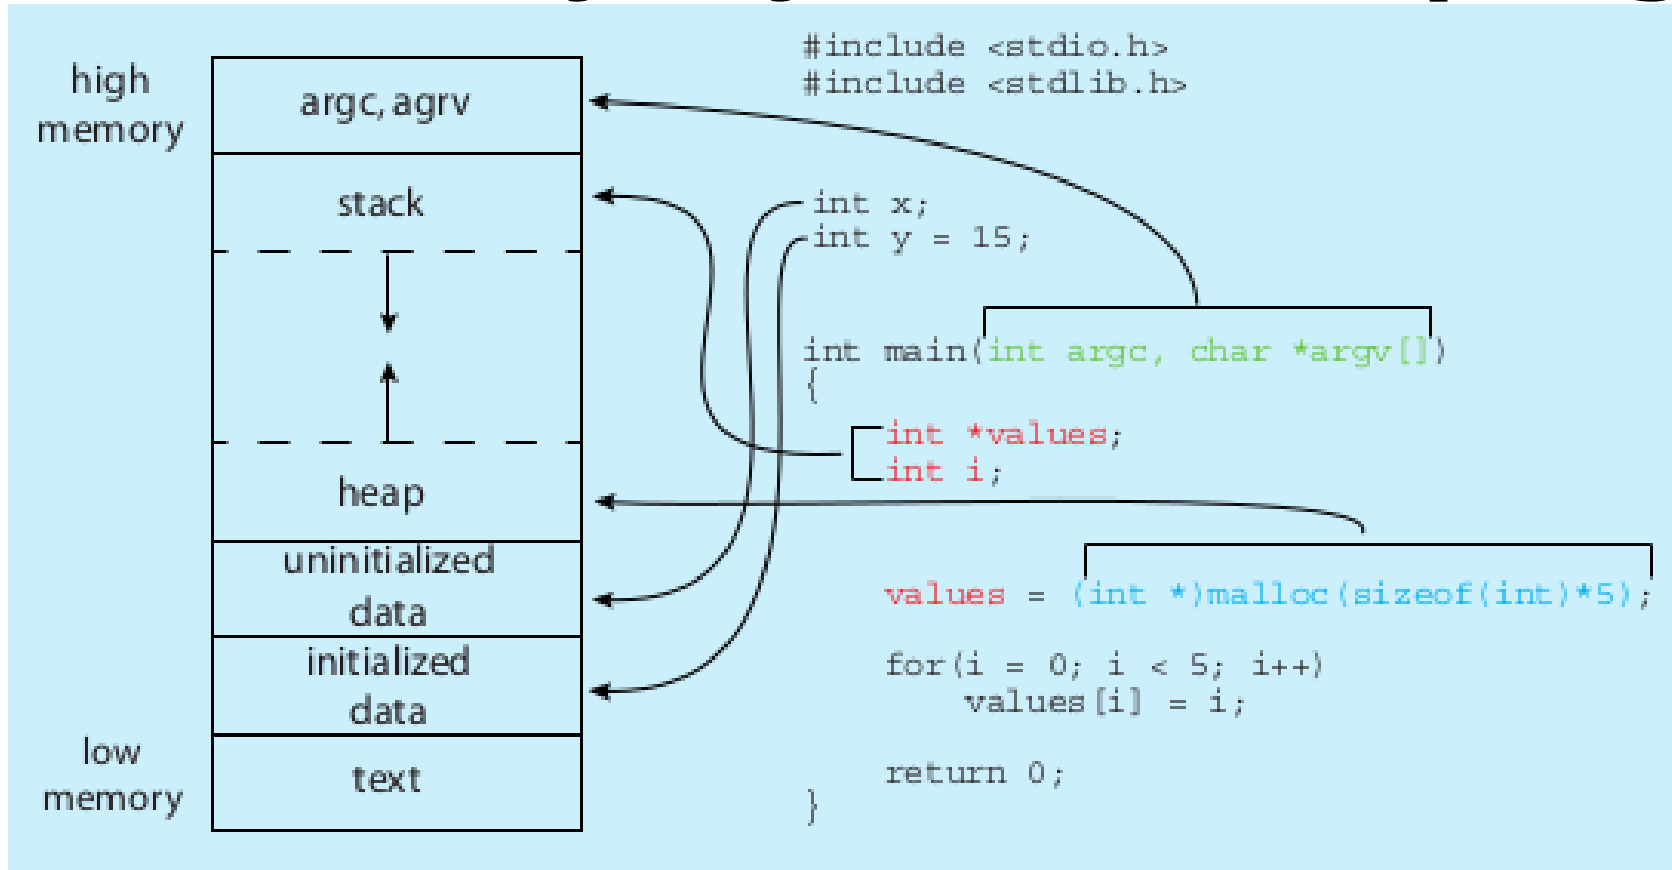
# Summary

- Understanding how the processor architecture drives the memory management features of OS and system programs (compilers, linkers)
- Understanding how different hardware designs lead to different memory management schemes by operating systems

# Addresses issued by CPU

- During the entire ‘on’ time of the CPU
  - Addresses are “issued” by the CPU on address bus
  - One address to fetch instruction from location specified by PC
  - Zero or more addresses depending on instruction
    - e.g. `mov $0x300, r1` # move contents of address 0x300 to r1 --  
> one extra address issued on address bus

# Memory layout of a C program



This “layout” shows (a) which parts of a C program occupy memory, when the program is running (b) A typical conceptual layout assumed by many compilers for calculating addresses in the generated machine code. This does not mean that other layouts are not possible

**\$ size /bin/ls**

text	data	bss	dec	hex	filename
128069	4688	4824	137581	2196d	/bin/ls

# Terminology

- Text
  - The machine code of the program : machine code for functions
- Data
  - Initiaized global variables
  - Do not get confused with the generic word “data” in English
- BSS
  - Uninitialized global variables
- Heap
  - Region from where malloc() gets memory
- Stack
  - Region from where the local variables and formal paramters are allocated space

# Desired from a multi-tasking system

- Multiple processes in RAM at the same time (multi-programming)
- Processes should not be able to see/touch each other's code, data (globals), stack, heap, etc.
- Further advanced requirements
  - Process could reside anywhere in RAM
  - Process need not be continuous in RAM
  - Parts of process could be moved anywhere in RAM

# Different 'times'

- Different actions related to memory management for a program are taken at different times. So let's know the different 'times'
- Compile time
  - When compiler is compiling your C code
- Load time
  - When you execute `./myprogram` and it's getting loaded in RAM by loader i.e. `exec()`
- Run time
  - When the process is alive, and getting scheduled by the OS

# The sequence

- Do not forget this
  - Machine code is typically generated by compiler
  - This machine code is put in RAM by the Loader (part of kernel) , that is `exec()` , when it's requested to run that program
  - The CPU's PIPELINE will issue addresses on address bus as seen in the executable file
- So question arises
  - How does compiler put in addresses (for code, data, bss, local-variables, etc) in the executable file ?




# Different types of Address binding

- Compile time address binding
  - Address of code/variables is fixed by compiler
  - Very rigid scheme
  - Location of process in RAM can not be changed ! Non-relocatable code.
- Load time address binding
  - Address of code/variables is fixed by loader
  - Location of process in RAM is decided at load time, but can't be changed later
  - Flexible scheme, relocatable code
- Run time address binding
  - Address of code/variables is fixed at the time of executing the code
  - Very flexible scheme , highly relocatable code
  - Location of process in RAM is decided at load time, but CAN be changed later also

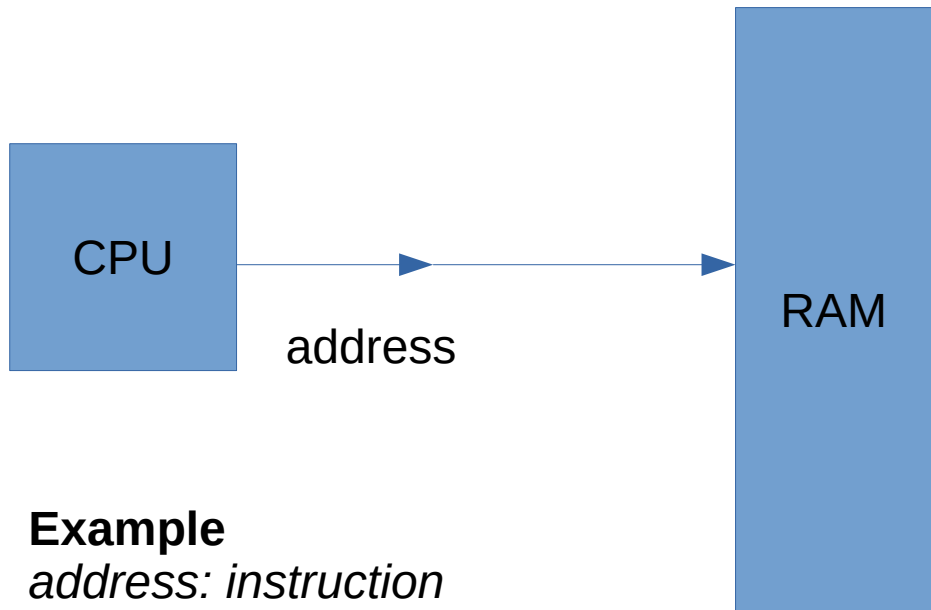
Which binding is actually used, is mandated by processor features + OS

# Simplest case



- Suppose the address issued by CPU reaches the RAM controller directly

# Simplest case



## Example


*address: instruction*

```
1000: mov $0x300, r1
1004: add r1, -3
1008: jnz 1000
```

Sequence of addresses sent by CPU: 1000, 0x300, 1004, 1008, 1000, 0x300, ...

- How does this impact the compiler and OS ?
- When a process is running the addresses issued by it, will reach the RAM directly
- So exact addresses of globals, addresses in “jmp” and “call” must be part the machine instructions generated by compiler
  - How will the compiler know the addresses, at “compile time” ?

# Simplest case




## Example

```
1000: mov $0x300, r1
1004: add r1, -3
1008: jnz 1000
```

- Solution: compiler assumes some fixed addresses for globals, code, etc.
- OS loads the program exactly at the same addresses specified in the executable file. **Non-relocatable code.**
- Now program can execute properly.

# Simplest case




## Example

```
1000: mov $0x300, r1
1004: add r1, -3
1008: jnz 1000
```

- Problem with this solution
  - Programs once loaded in RAM must stay there, can't be moved
  - What about 2 programs?
    - Compilers being "programs", will make same assumptions and are likely to generate same/overlapping addresses for two different programs
    - Hence only one program can be in memory at a time !
    - No need to check for any memory boundary violations – all memory belongs to one process
- Example: DOS

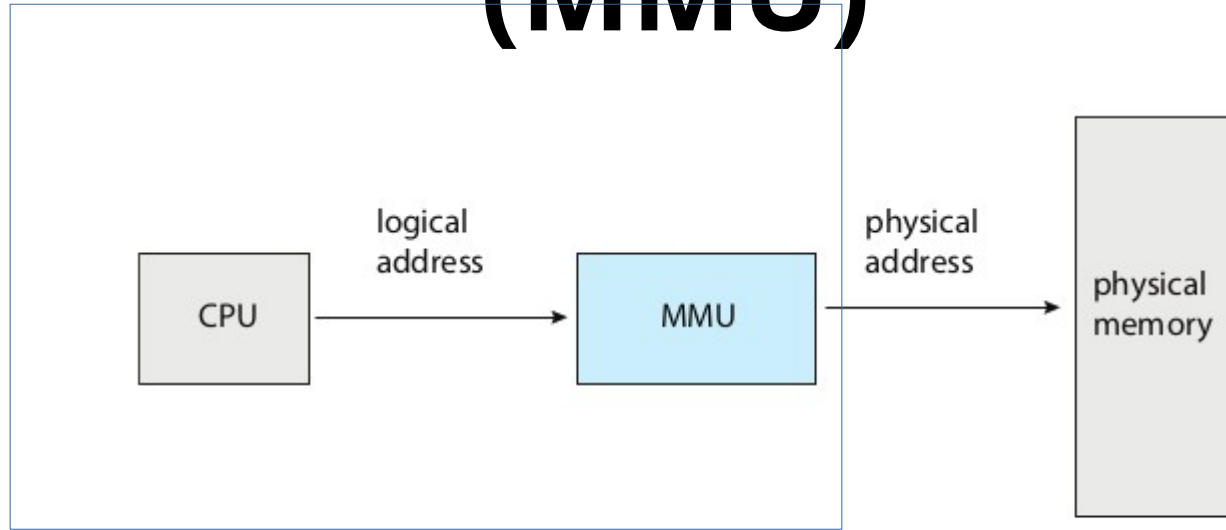
# Base/Relocation + Limit scheme



**Figure 9.6** Hardware support for relocation and limit registers.

- Base and Limit are two registers inside CPU's Memory Management Unit
- 'base' is added to the address generated by CPU
- The result is compared with base+limit and if less passed to memory, else hardware interrupt is raised

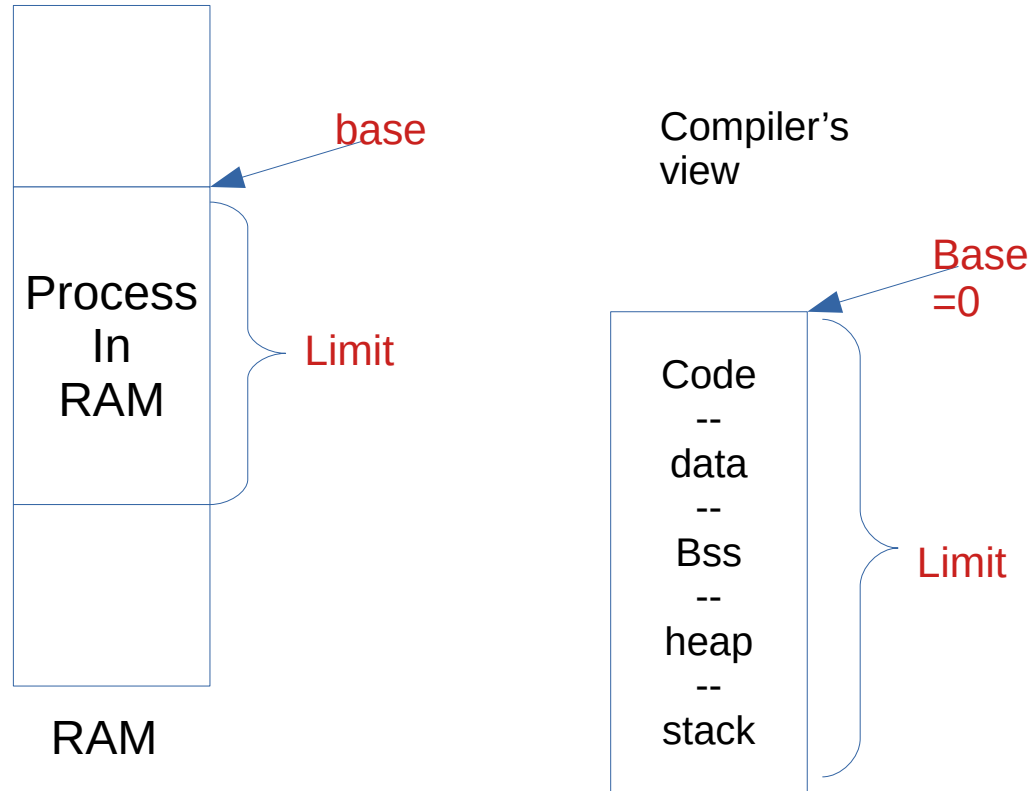
# Memory Management Unit (MMU)



**Figure 9.4** Memory management unit (MMU).

- Is part of the CPU chip, acts on every memory address issue by execution unit of the CPU
- In the scheme just discussed, the base, limit calculation parts are part of MMU

# Base/Relocation + Limit scheme



- **Compiler's work**
  - Assume that the process is one continuous chunk in memory, with a size limit
  - Assume that the process starts at address zero (!) and calculate addresses for globals, code, etc. And accordingly generate machine code



# Base/Relocation + Limit scheme

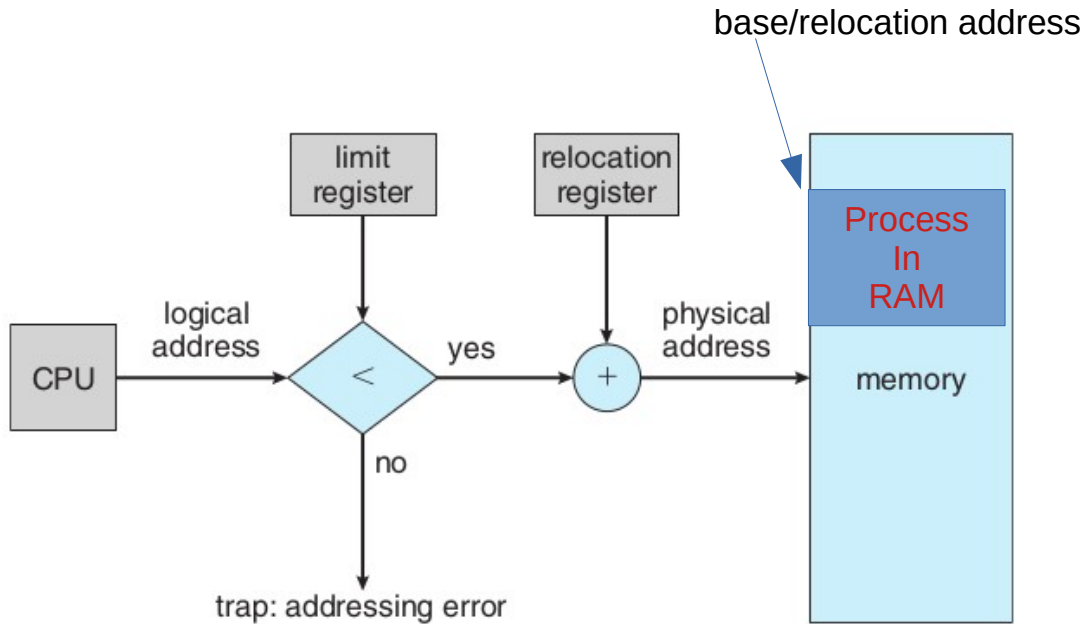


Figure 9.6 Hardware support for relocation and limit registers.

- **Loader's job**
  - While loading the process in memory – must load as one continuous segment
    - Find an empty slot for this!
  - Remember the 'base' and 'limit' values in the struct proc (memory management info in PCB)
    - Setup the limit to be the size of the process as set by compiler in the executable file.
  - Call the scheduler()
- **Scheduler()**
  - Will replace the base-limit registers with values obtained from PCB before scheduling the process

# Base/Relocation + Limit scheme

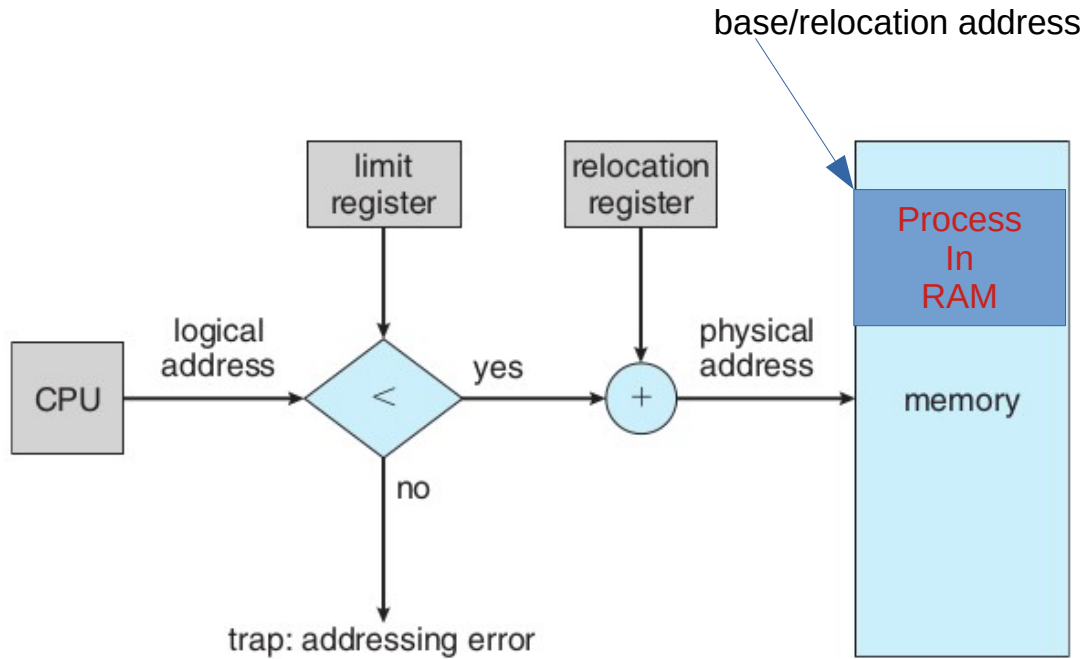
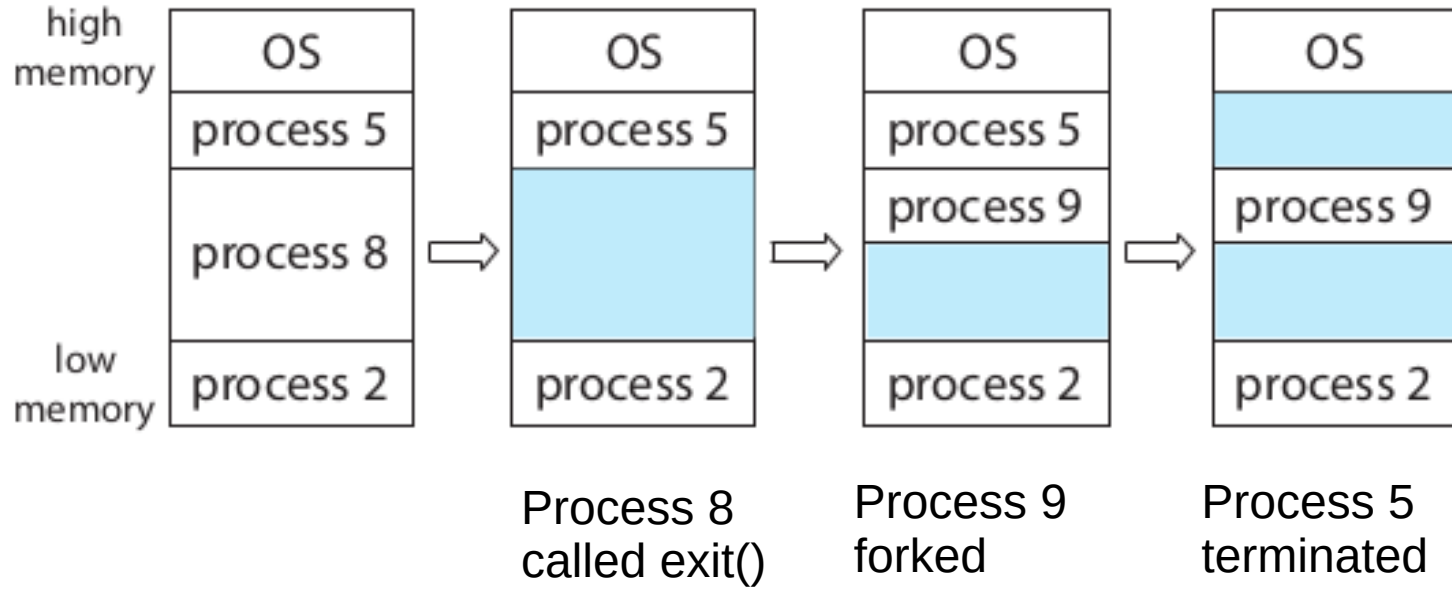



Figure 9.6 Hardware support for relocation and limit registers.

- Combined effect
  - **“Relocatable code”** – the process can go anywhere in RAM at the time of loading
  - Some memory violations can be detected – a memory access beyond base+limit will raise interrupt, thus running OS in turn, which may take action against the process

# Example scenario of memory in base+limit scheme



# Continuous memory management and external fragmentation problem



Free chunks: 2 MB, 1MB, 1MB, 1MB

Total 5 MB is available!

Can we create a process of size 3MB?

No! - 3 MB not continuous!

# External Fragmentation

- OS needs to find a continuous free chunk of memory that fits the size of the “segment”
  - If not available, your `exec()` can fail due to lack of memory
- Suppose 50k is needed
  - Possible that among 3 free chunks total 100K may be available, but no single chunk of 50k!
  - **External fragmentation**
- Solution to external fragmentation: **compaction** – move the chunks around and make a continuous big chunk available. Time consuming, tricky.

It should be possible to have relocatable code  
even with “simplest case”

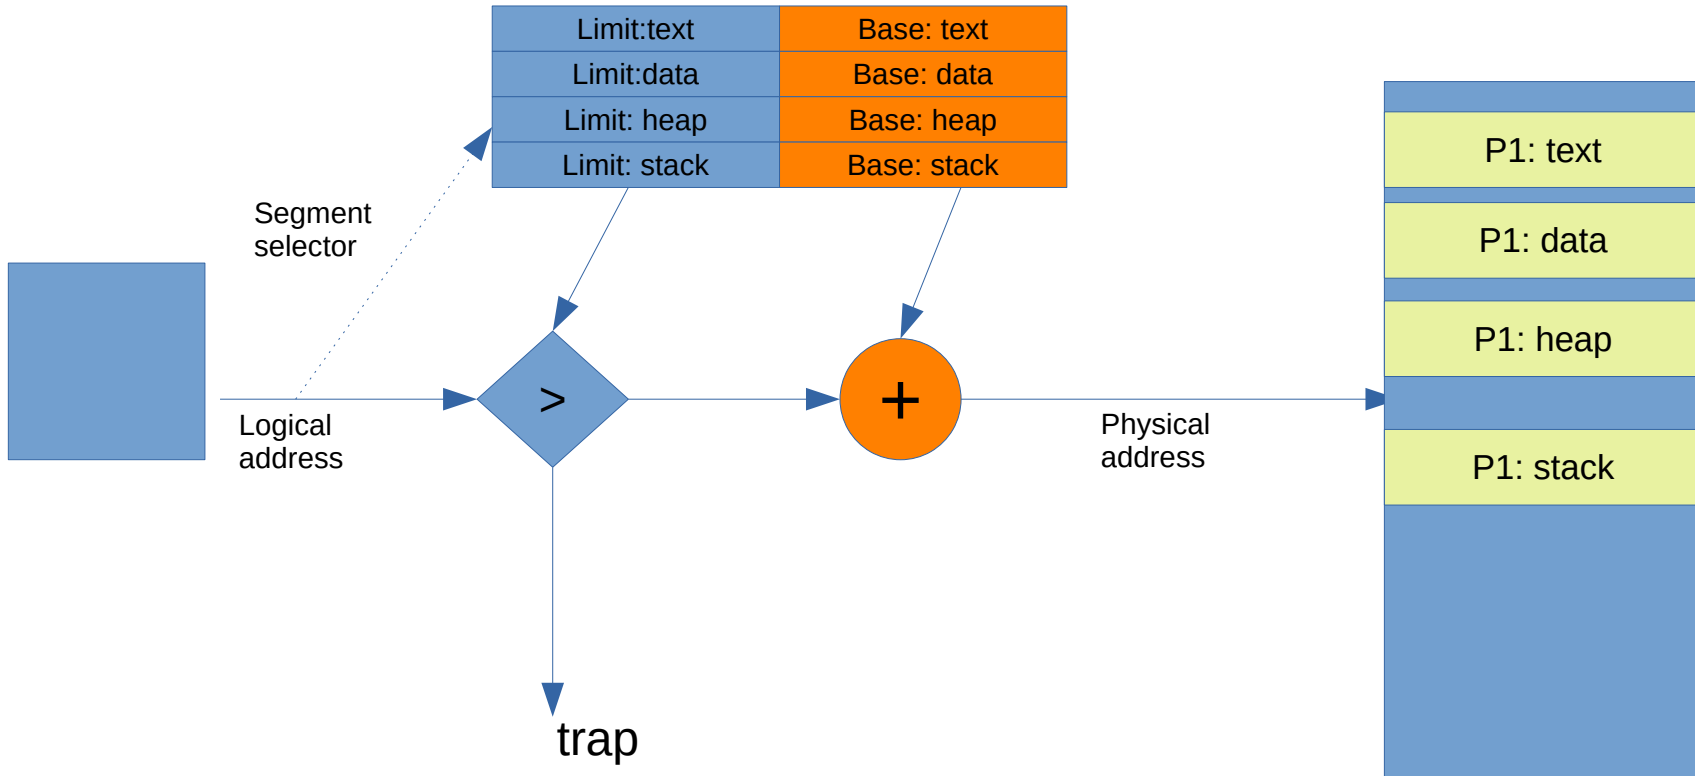
By doing extra work during “loading”.

How?

(Ans: loader replaces all addresses in the code!  
That is quite a lot of work, and challenging too)

# Next scheme:

## Multiple base +limit pairs



# Specifying “selector”

- Explicitly
  - `Mov ES: 300, $30`
  - Here selector is Extra Segment Register and Logical address (or offset) is 300
- Implicitly, like in x86
  - `Mov 300, $30`
  - Here Data Segment register is “implicit” selector



# Next scheme: Segmentation

## Multiple base +limit pairs

- Multiple sets of base + limit registers
- Whenever an address is issued by execution unit of CPU, it will also include reference to some base register
  - And hence limit register paired to that base register will be used for error checking
- Compiler: can assume a separate chunk of memory for code, data, stack, heap, etc. And accordingly calculate addresses . Each “segment” starting at address 0.
- OS (Loader): will load the different ‘sections’ in different memory regions and accordingly set different ‘base’ registers

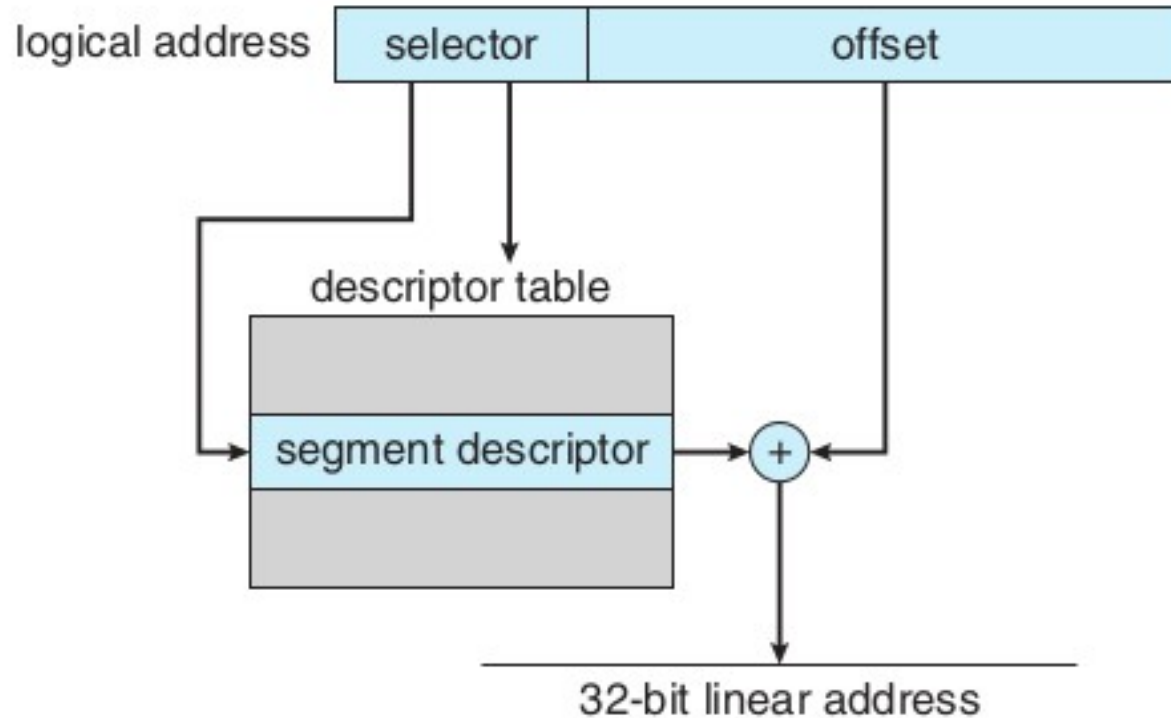
# Next scheme:

## Multiple base +limit pairs, with further indirection

- Base + limit pairs can also be stored in some memory location (not in registers). Then it's called **Segment Table**.
  - Question: how will the cpu know where it's in memory?
  - One CPU register to point to the location of table in memory . **Segment Table Base Register (STBR)**
    - X86 has two tables. Local Descriptor Table and Global Descriptor Table. Both are segment tables.
    - Accordingly it has LDTR and GDTR.
- Segment registers still in use, they give an index in this table
- This is x86 segmentation
  - Flexibility to have lot more “base+limits” in the array/table in memory

# Next scheme:

## Multiple base +limit pairs, with further indirection



Note:

"Selector" here is normally a segment register, like CS, DS

The machine code only has offset in it

The "Descriptor table" is a segmentation table

Figure 9.22 IA-32 segmentation.

# Segmentation and External fragmentation

- Does segmentation also suffer from external fragmentation?
  - Yes!

# How many segment tables?


- One per process
- One for the kernel!
  - Remember: kernel code also undergoes translation in MMU!

# A point of confusion on “root” user

- “root” user has privileges in accessing files
- When “root” user runs an application, it runs in process-mode, not in kernel-mode!
  - That process can do “more” but , it’s not kernel mode code!
- Kernel is not run by “root” user!
  - Kernel is not run by any user

# Solution to external fragmentation

- Compaction !
- OS moves the process chunks in memory to make available continuous memory region
  - Then it must update the memory management information in PCB (e.g. base of the process) of each process
- Time consuming
- Possible only if the relocation+limit scheme of MMU is available



# Another solution to external fragmentation: Fixed size partitions

- Fixed partition scheme
- Memory is divided by OS into chunks of equal size:  
e.g., say, 50k
  - If total 1M memory, then 20 such chunks
- Allocate one or more chunks to a process, such that the total size is  $\geq$  the size of the process
  - E.g. if request is 50k, allocate 1 chunk
  - If request is 40k, still allocate 1 chunk
  - If request is 60k, then allocate 2 chunks
- Leads to internal fragmentation
  - space wasted in the case of 40k or 60k requests above


50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K
50K

Kernel
Kernel
Free
P1 (50 KB)
P2 (80 KB)
Unused (20 KB)
Free
P3 (120 KB)
Unused (30 KB)
Free
Free



# Solving external fragmentation problem

- Process should not be continuous in memory!
  - The trouble is finding a big continuous chunk!
- Divide the continuous process image in smaller chunks (let's say 4k each) and locate the chunks anywhere in the physical memory
  - Need a way to map the *logical* memory addresses into *actual physical memory addresses*



**Figure 9.9** Paging model of logical and physical memory.

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

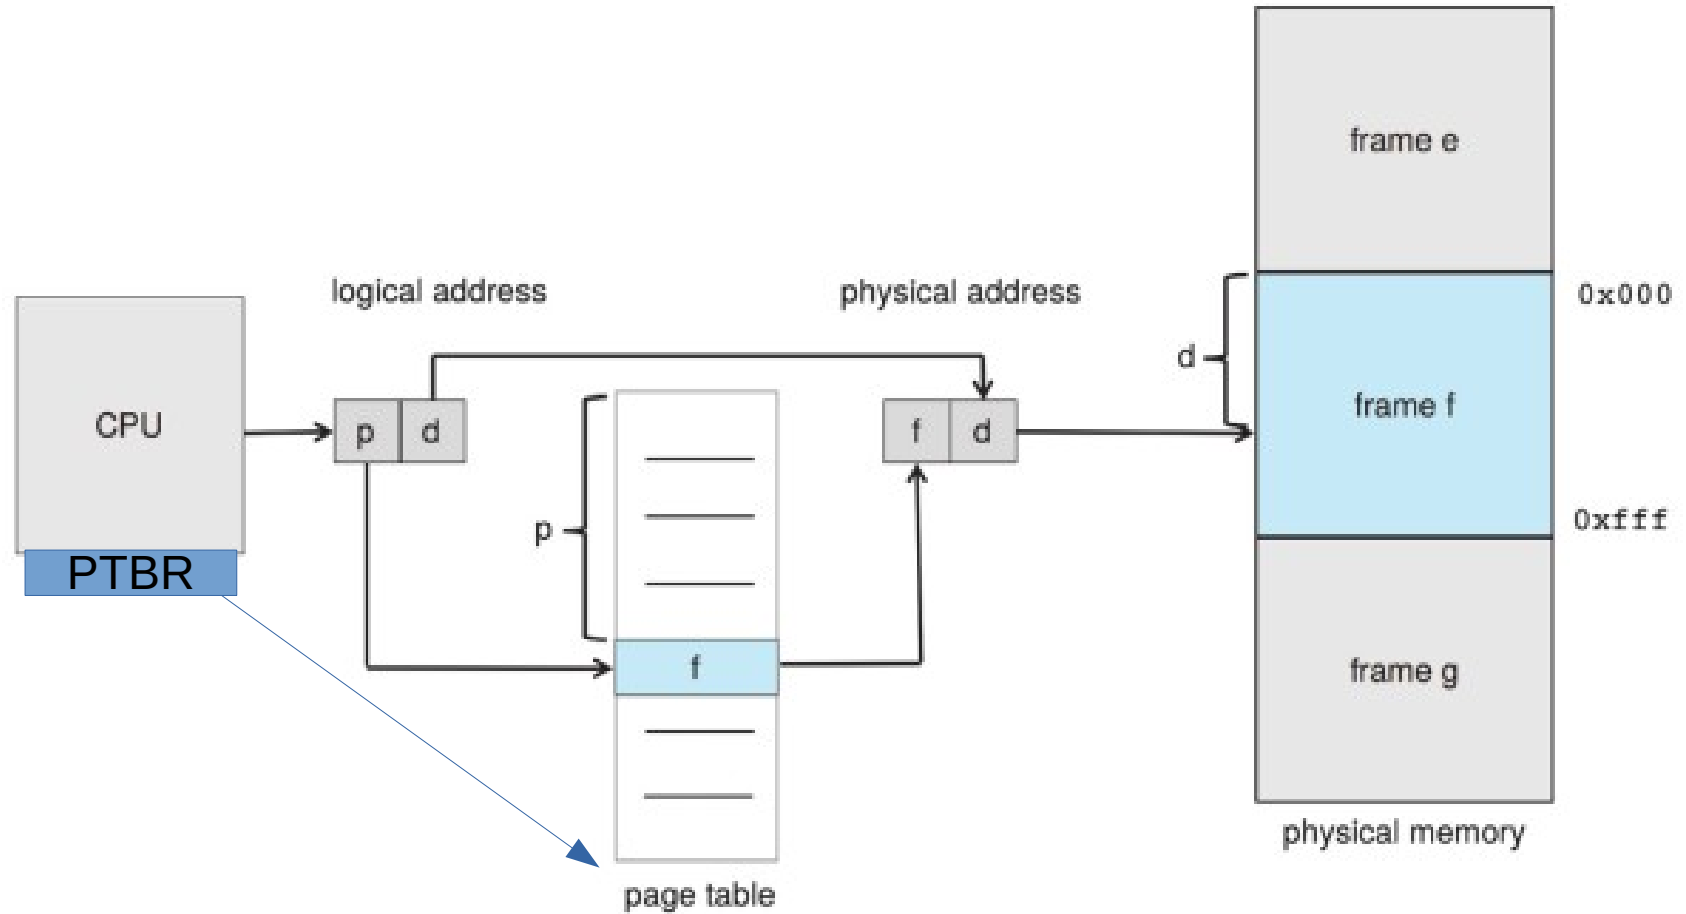
0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

**Figure 9.10** Paging example for a 32-byte memory with 4-byte pages.




**Figure 9.8** Paging hardware.

# Paging

- Process is assumed to be composed of equally sized “pages” (e.g. 4k page)
- Actual memory is considered to be divided into page “frames”.
- CPU generated logical address is split into a page number and offset
- A Page Table Base Register (PTBR) inside CPU will give location of an in memory table called page table
- Page number used as offset in a table called page table, which gives the physical page frame number
- Frame number + offset are combined to get physical memory address

# Paging


- Compiler: assume the process to be one continuous chunk of memory (!) . Generate addresses accordingly . Compiler still treats text,data,bss,.. to be separate chunks, but in multiples of page-sizes.
- OS: at exec() time – allocate different frames to process, allocate a page table(!), setup the page table to map page numbers with frame numbers, setup the page table base register, start the process
- Now hardware will take care of all translations of logical addresses to physical addresses



**Figure 9.11** Free frames (a) before allocation and (b) after allocation.

# Speeding up paging

- Translation Lookaside Buffer (TLB)
- Part of CPU hardware
- A cache of Page table entries
- Searched in parallel for a page number





12,287  
10,468

page 5
page 4
page 3
page 2
page 1
page 0

00000

frame number      valid-invalid bit


0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page n

# Memory protection with paging

**Figure 9.13** Valid (v) or invalid (i) bit in a page table.



**Shared  
pages (e.g.  
library)  
with paging**

**Figure 9.14** Sharing of standard C library in a paging environment.



# Paging: problem of large PT

- 64 bit address
- Suppose 20 bit offset
  - That means  $2^{20} = 1 \text{ MB}$  pages
  - 44 bit page number:  $2^{44}$  that is trillion sized page table!
  - Can't have that big continuous page table!

# Paging: problem of large PT

- 32 bit address
- Suppose 12 bit offset
  - That means  $2^{12} = 4 \text{ KB}$  pages
  - 20 bit page number:  $2^{20}$  that is a million entries
  - Can't always have that big continuous page table as well, for each process!

# Hierarchical paging



**Figure 9.15** A two-level page-table scheme.

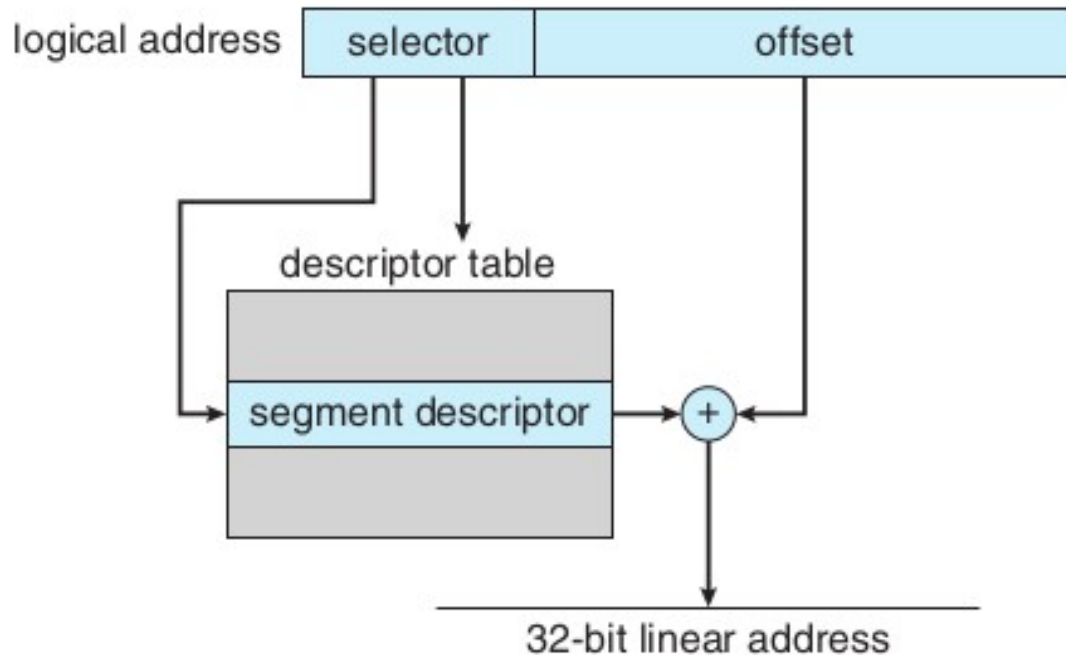
outer page	inner page	offset
$p_1$	$p_2$	$d$
42	10	12

# X86 memory management



**Figure 9.21** Logical to physical address translation in IA-32.

# Segmentation in x86

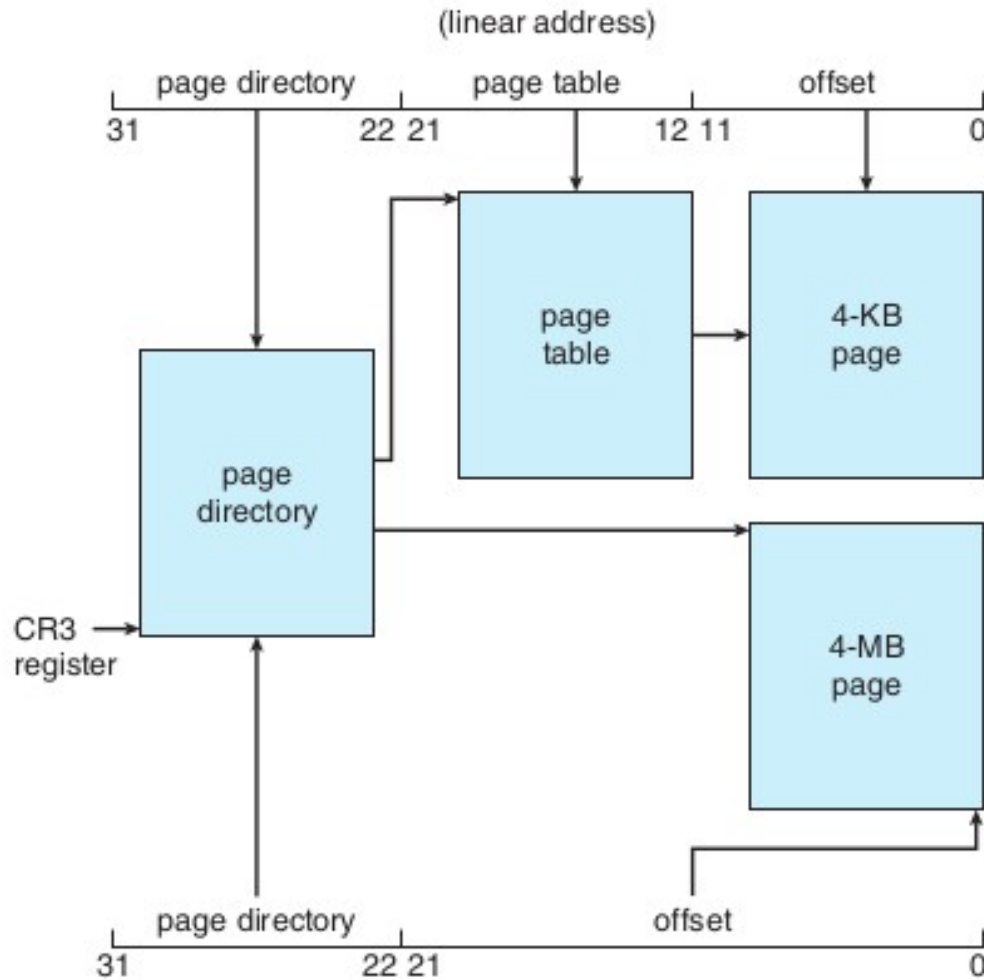


- The selector is automatically chosen using Code Segment (CS) register, or Data Segment (DS) register depending on which type of memory address is being fetched
- Descriptor table (that is a segmentation table) is in memory
- The location of Descriptor table (Global DT- GDT or Local DT - LDT) is given by a GDT-register i.e. GDTR or LDT-register i.e. LDTR

**Figure 9.22** IA-32 segmentation.



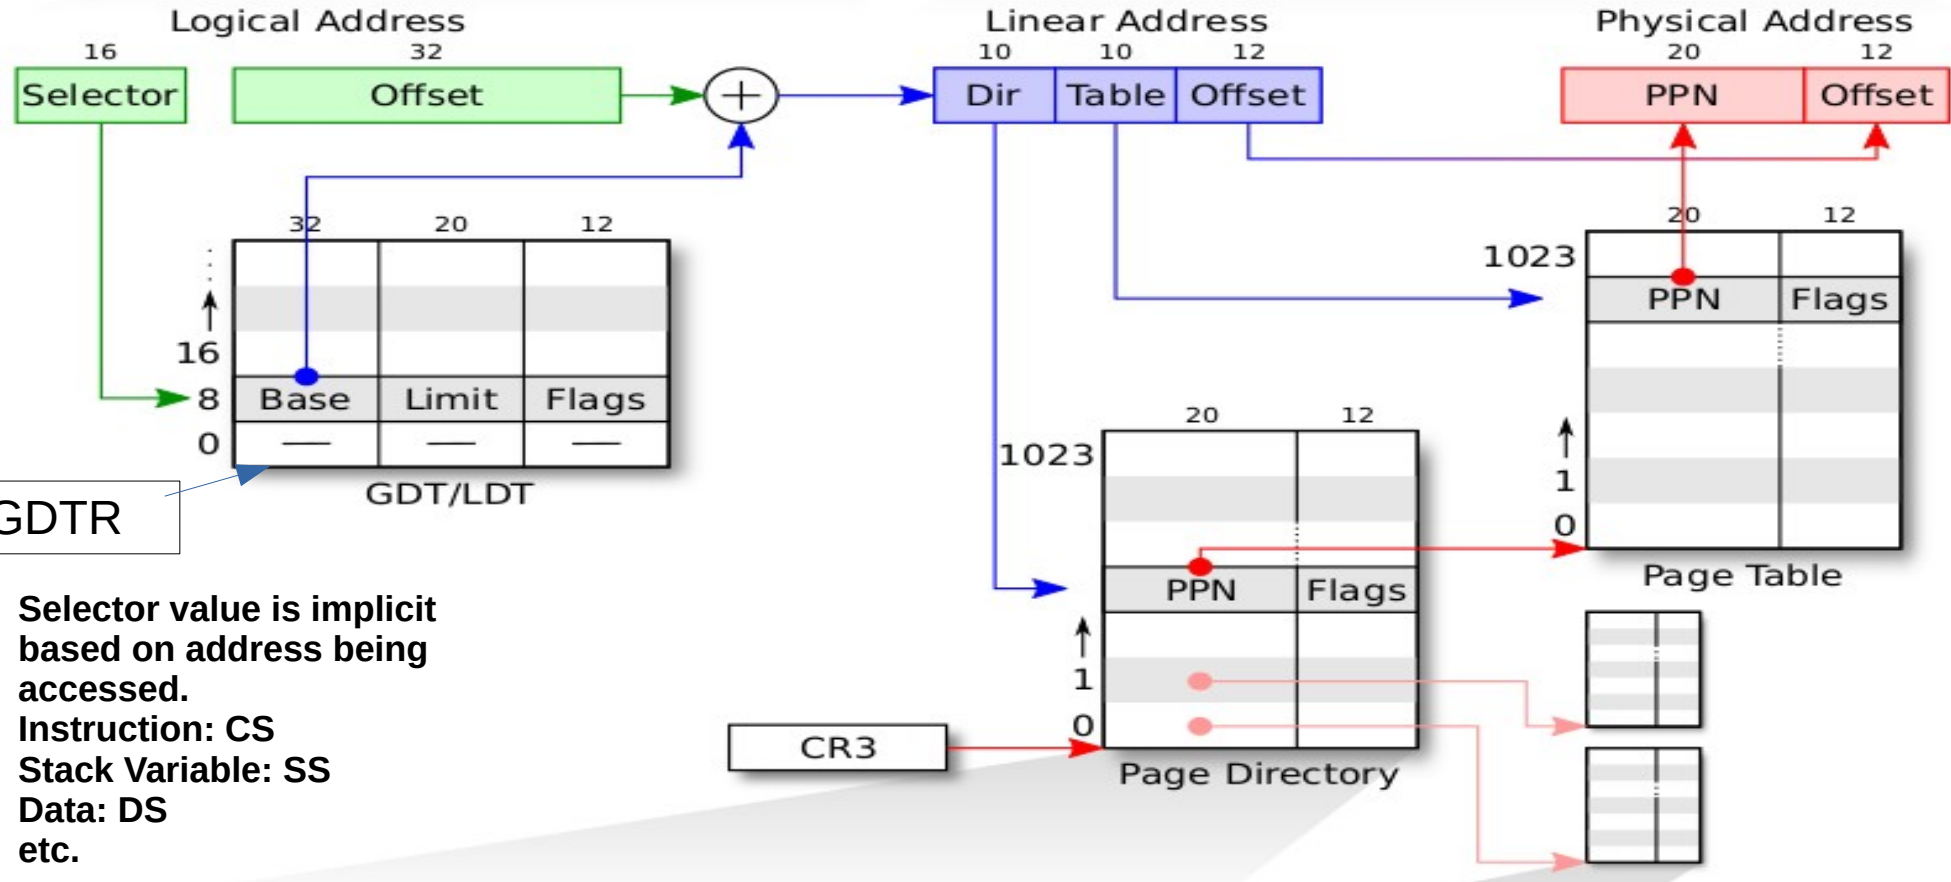
# Paging in x86



**Figure 9.23** Paging in the IA-32 architecture.

- Depending on a flag setup in CR3 register, either 4 MB or 4 KB pages can be enabled
- Page directory, page table are both in memory

# X86 Segmentation + Paging



# X86 Segmentation + Paging

- **Paging is optional, segmentation compulsory**
  - **Setting flags in Control Registers (CR) enables this**
- **Page Table, Page Directory, page - are all size=4k, if 2-level paging is used**
  - **Makes life simpler for the kernel**