

# **Chapter-I and II**

## **Introduction to OS**

Abhijit A. M.  
abhijit.comp@coep.ac.in

# OS or kernel?


- **Debate on terminology**
- **What you use in daily life is not kernel, but**
  - GUI, Shell, Applications, ...
  - One view: OS = kernel + (GUI, Shell, Libraries, System programs, Minimum Applications,...)
  - Correct, IMO!
- **What we study in this course is kernel**
  - Not how is “Ubuntu” built!

# Building an “OS”

- **E.g. Debian**
  - A collection of thousands of applications, libraries, system programs, ... and Linux kernel !
  - Linux kernel is at the heart, but heart is not a human without the body !
  - Job of “Debian Developers”
    - Collect the source code of all things you want
    - Create an “Environment” for compiling things : bootstrap challenge
    - Compile everything
    - Ensure that things work “with each other”
      - Very difficult! Versions, dependencies !
    - “Package” things (e.g. make .deb files )

# Components of a computer system

## Abstract view

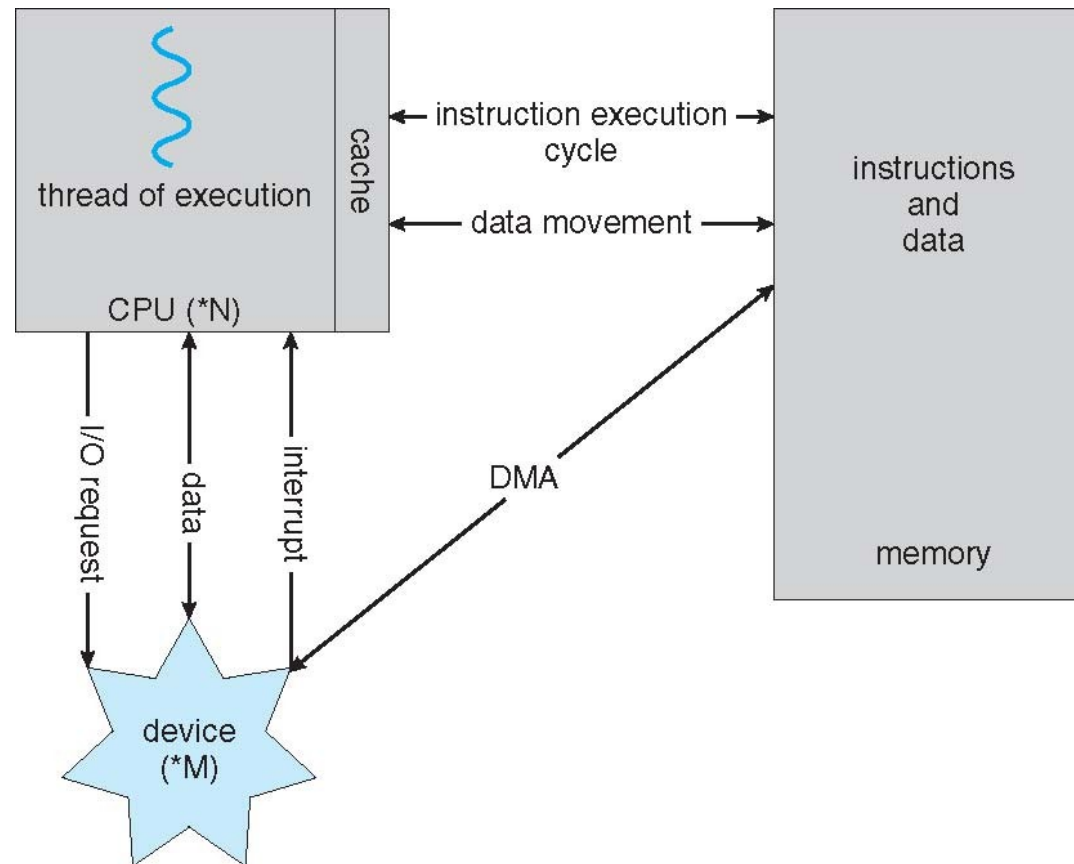


**Figure 1.1** Abstract view of the components of a computer system.

# Today, different type of kernels to serve different needs

- Each “Computing Environment” has a different requirement of managing resources, processes, etc
  - Desktop / Laptop
  - Thin Clients
  - Workstations
  - Handheld
  - “Servers”
  - Distributed systems
  - Peer to peer systems
  - Virtualization
  - Cloud
  - Real time

# Von Neumann architecture




*A von Neumann architecture*

# **A key to “understanding”**

- **Everything happens on processor**
- **Processor is always running some instruction**
- **We should be able to tell possible execution sequences on processor**

# Memory hierarchy





# Memory hierarchy

Level	1	2	3	4	5
Name	registers	cache	main memory	solid state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25 - 0.5	0.5 - 25	80 - 250	25,000 - 50,000	5,000,000
Bandwidth (MB/sec)	20,000 - 100,000	5,000 - 10,000	1,000 - 5,000	500	20 - 150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

# System calls

- **We have seen**
  - **Fork(), exec(): system calls to create processes**
  - **Open(), read(), write() : system calls to play with files**
- **System call: A service given by kernel**

# System calls: the problem

- **On a time shared system**
  - Applications and kernel run on the CPU taking turns
  - Kernel allocates time to application , by setting a value in timer register, and then makes application run
    - **Timer ticks down every clock cycle**
  - *Timer hardware interrupt* occurs when timer = 0 , that is time of application is over
  - Kernel runs again

# System calls: the problem

- **Kernel must provide “services” to applications**
  - Create processes
  - Grant access to resources to processes
  - Etc
- **Process should be able to “call” kernel’s service**
- **But everything runs on processor!**
  - That is both processes and kernel
- **Kernel should be able to run particular instructions which processes should not be able to !**

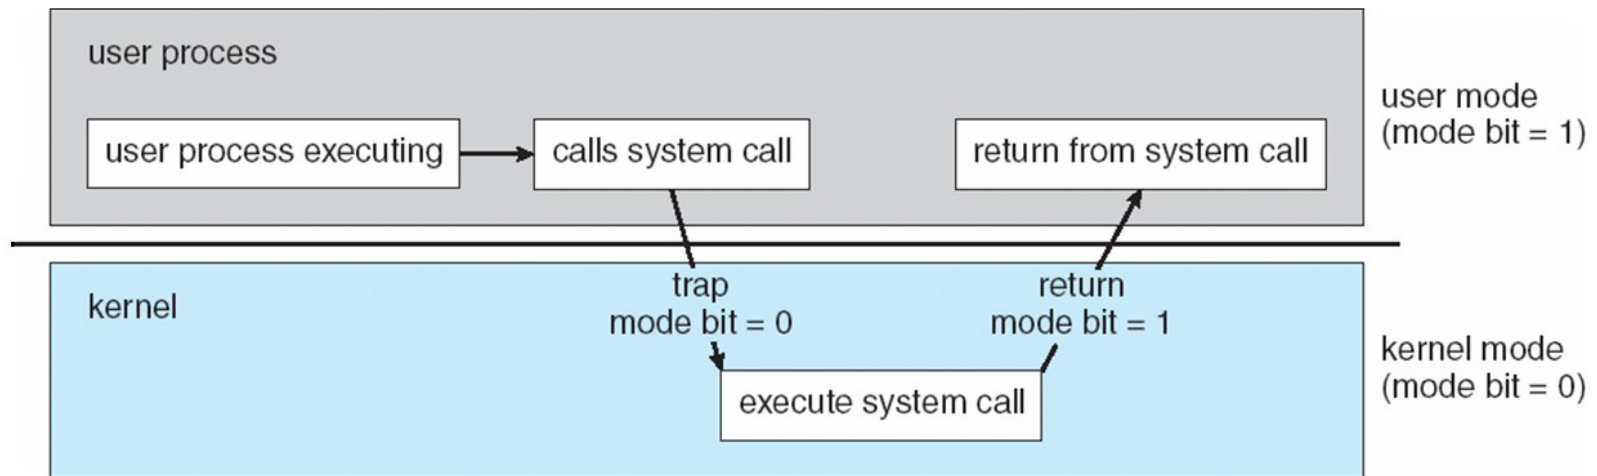
# System calls: the problem

- **Privileged Instructions**
  - Set the timer register
  - Set a value in registers of I/O devices
  - Set the MMU
  - Etc
- **How to ensure that only kernel code can run privileged instructions?**
  - Not possible without hardware support in CPU!

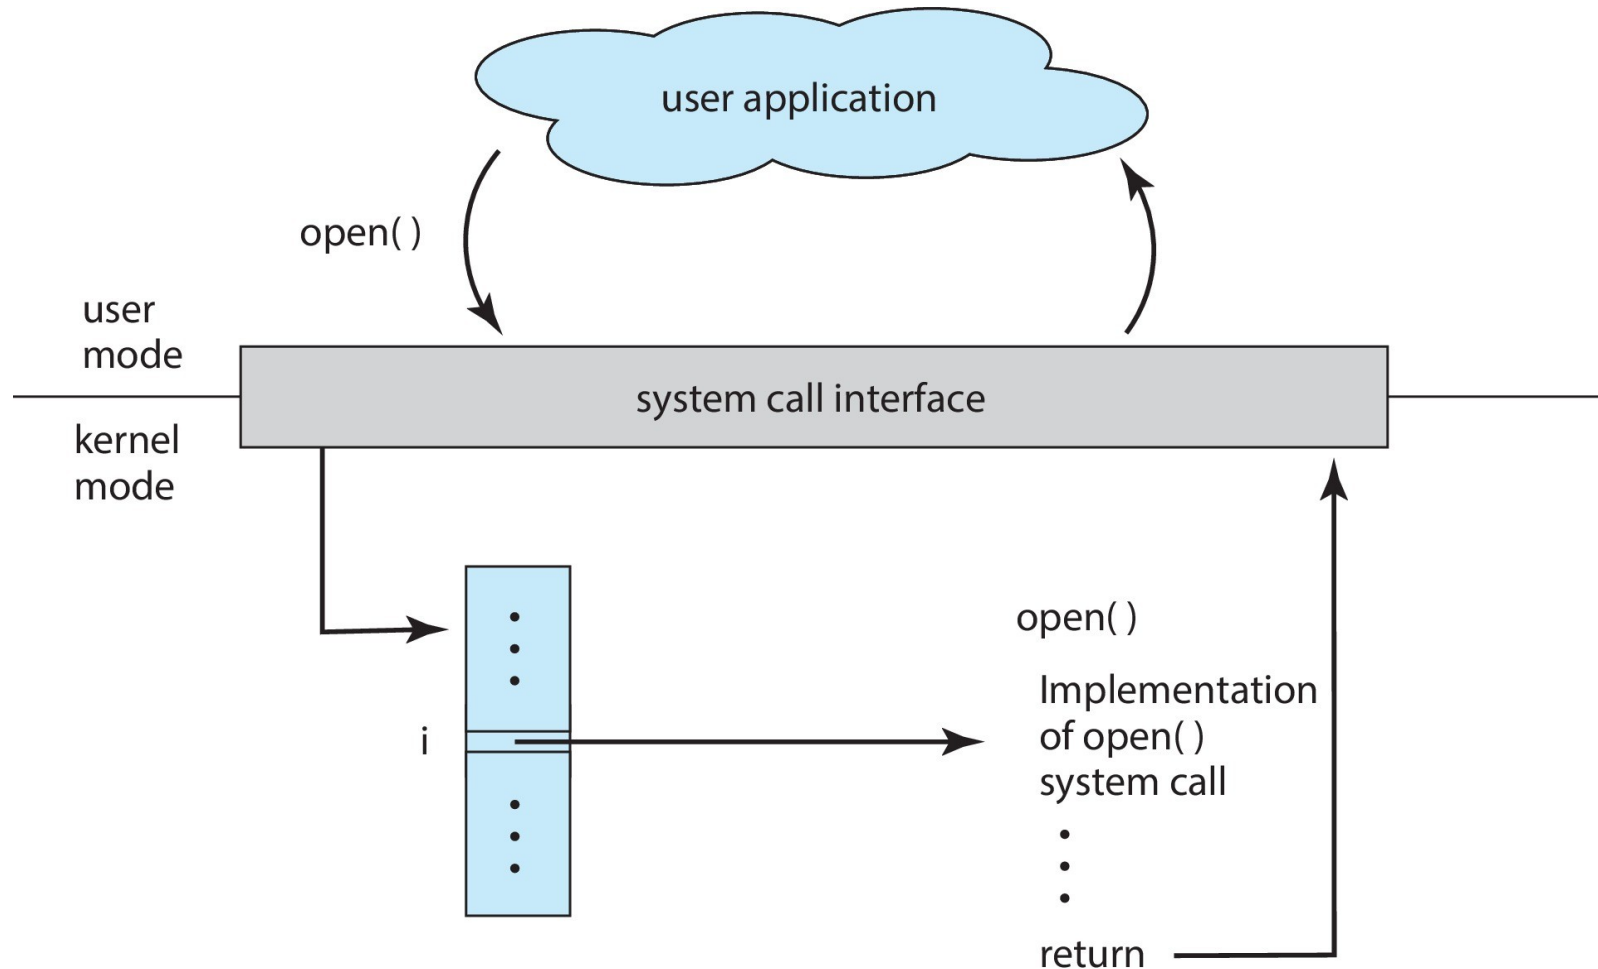
# System calls: the problem

- **CPU: 2 modes**
  - User mode and kernel mode
  - All instructions can be executed in kernel mode
  - Non-privileged instructions can be executed in User mode
  - How to ensure this?
- **Special instruction: software interrupt instruction**
  - Changes the mode from User Mode to Kernel mode and changes the CS+IP (or PC) to a pre-defined address (where the kernel already exists)
  - Two things at a time – that's the key

# 2 modes of operation of CPU needed for multi-tasking system

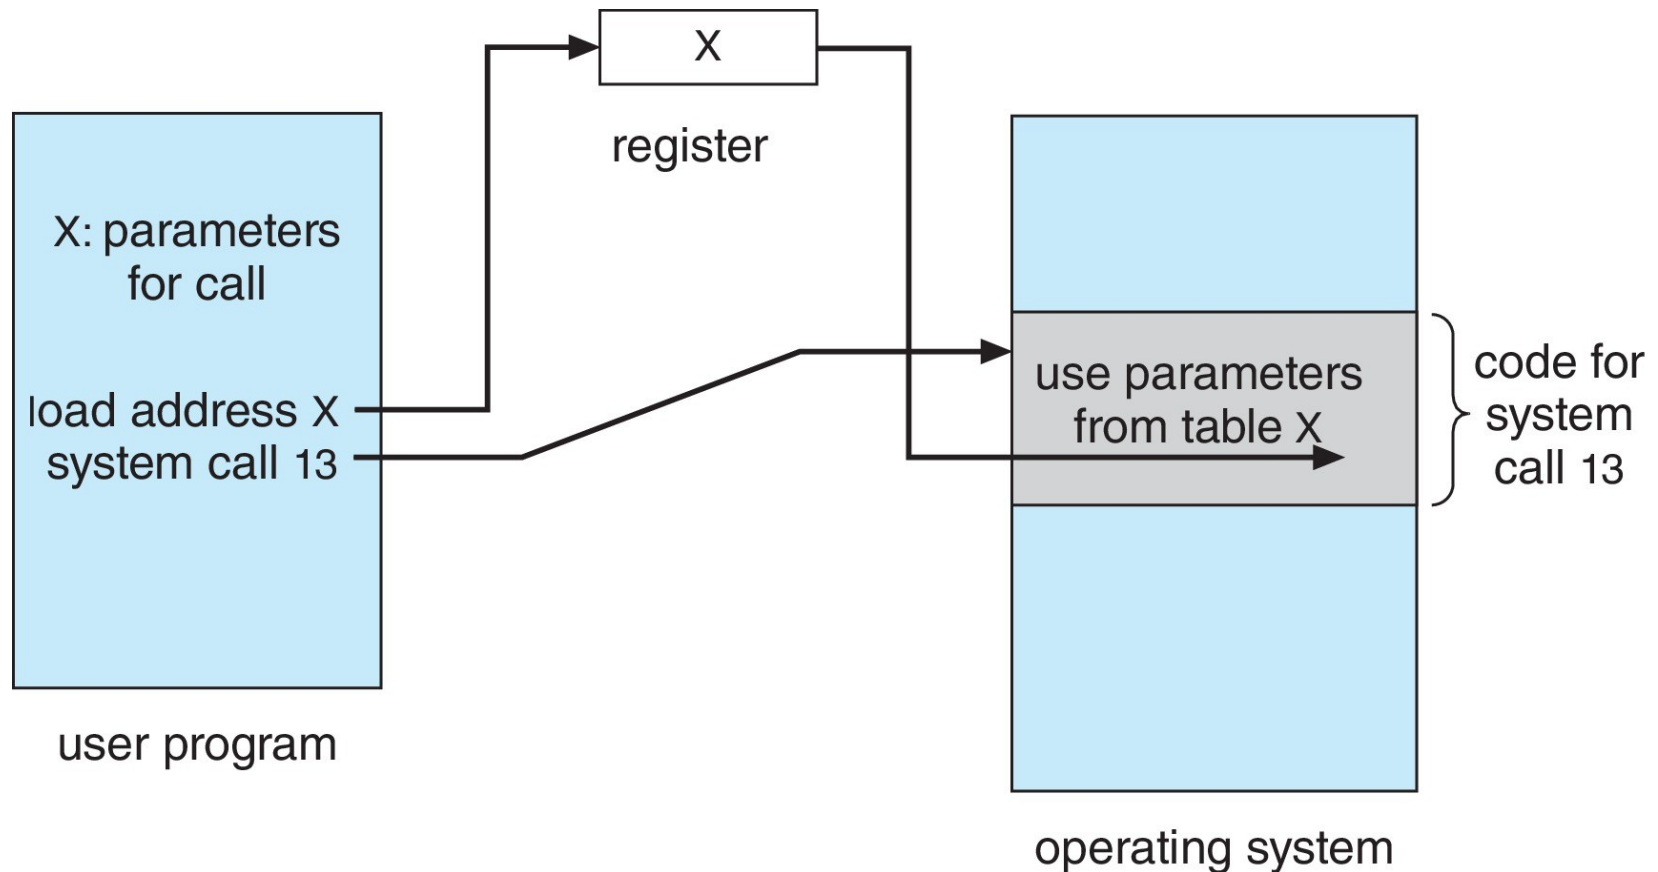


# Application – kernel interaction in system calls





# System call – parameter passing



Sharing of data between “user mode” and “kernel mode” code !

# What a modern multi-processor, multi-tasking kernel does

- **Process Management**

- Create an environment (fork, exec, exit, etc) so that processes can be created, exited, ...

- **Resource management**

- Processes can access “resources” only through kernel
- System calls like: open(), brk(), ...
- Includes storage management – enable a “tree” type view of files and folders
- Manage memory for itself and all processes
- I/O Management – make I/o devices accessible to processes and itself
- Manage users (uid, gid, setuid(), seteuid(), .. )

- **Protection and Security**

- Limit what a process can do
- Prevent “attacks” by internal and external agents

# System calls exist for ...

- **Process control**
  - create process (fork), terminate process (exit)
  - end, abort (kill)
  - load, execute (exec)
  - get process attributes (getpid, ..), set process attributes (setrlimit, ...)
  - wait for time (sleep, ...)
  - wait event (wait, ...), signal event (signal, kill, ... )
  - allocate and free memory (brk, ... )
  - Dump memory if error
  - Debugger for determining bugs, single step execution
  - Locks for managing access to shared data between processes

# System calls exist for ...

- **File management**

- create file (open, create), delete file (unlink)
- open, close (close()) file
- read, write, reposition (read, write, lseek, ...)
- get and set file attributes (chown, chmod, stat, ... )

- **Device management**

- request device, release device
- read, write, reposition
- get device attributes, set device attributes
- logically attach or detach devices
- On Linux: mostly done using file syscalls, but special calls are also there: mknod, mkfifo, ...

# System calls exist for ...

- **Information maintenance**
  - get time or date, set time or date
  - get system data, set system data
  - get and set process, file, or device attributes
- time, fcntl, ...
- **Communications**
  - create, delete communication connection
  - send, receive messages if message passing model to host name or process name
  - From client to server
  - Shared-memory model create and gain access to memory regions
  - transfer status information
  - attach and detach remote devices
- Examples: socket, bind, listen, shmget, send, recv,

# System calls exist for ...

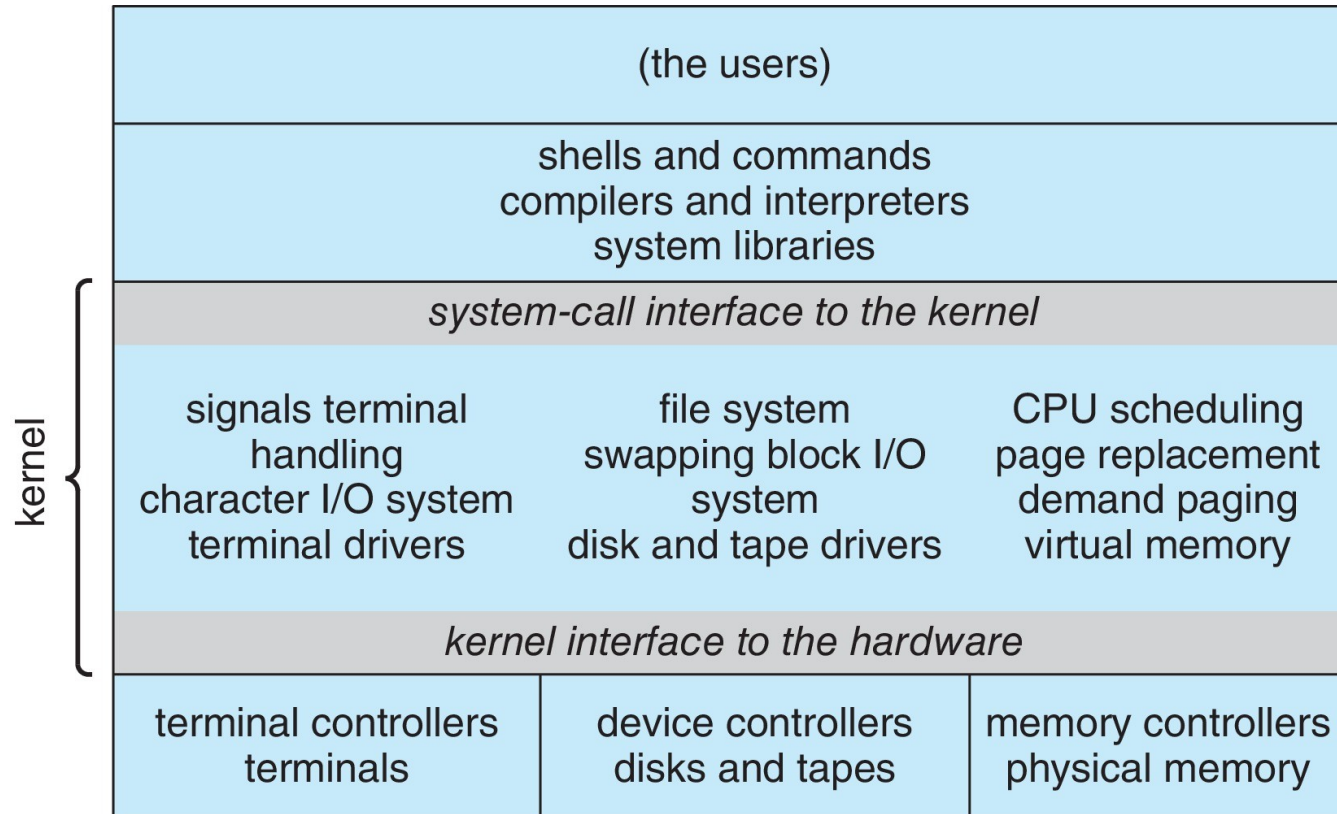
- **Protection**

- **Control access to resources**
- **Get and set permissions**
- **Allow and deny user access**

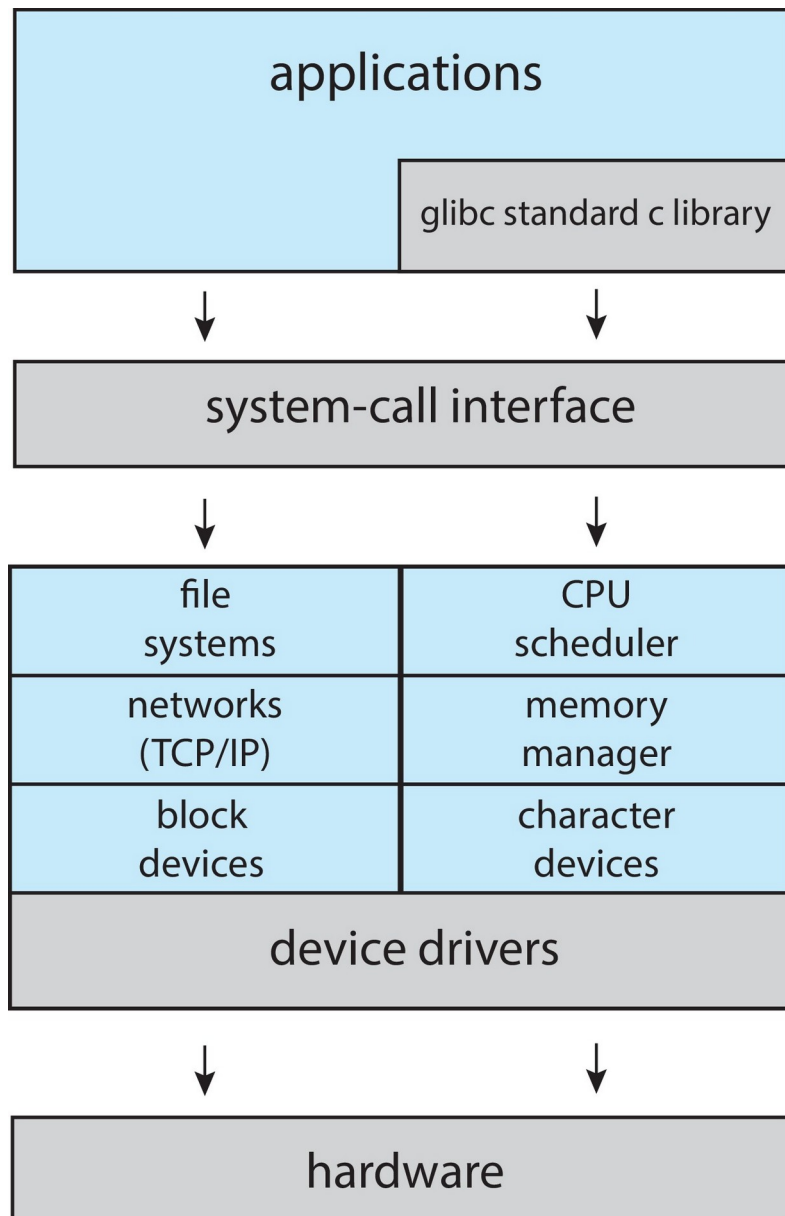
- **More**

- **On LinuxL “man syscalls”**
- **POSIX: Portable Operating System Interface**
- **See output of strace**

# Traditional Unix system structure



# Linux system structure




- **Linux kernel is extendible**
  - **Kernel modules**
  - **lsmod, insmod, ..**



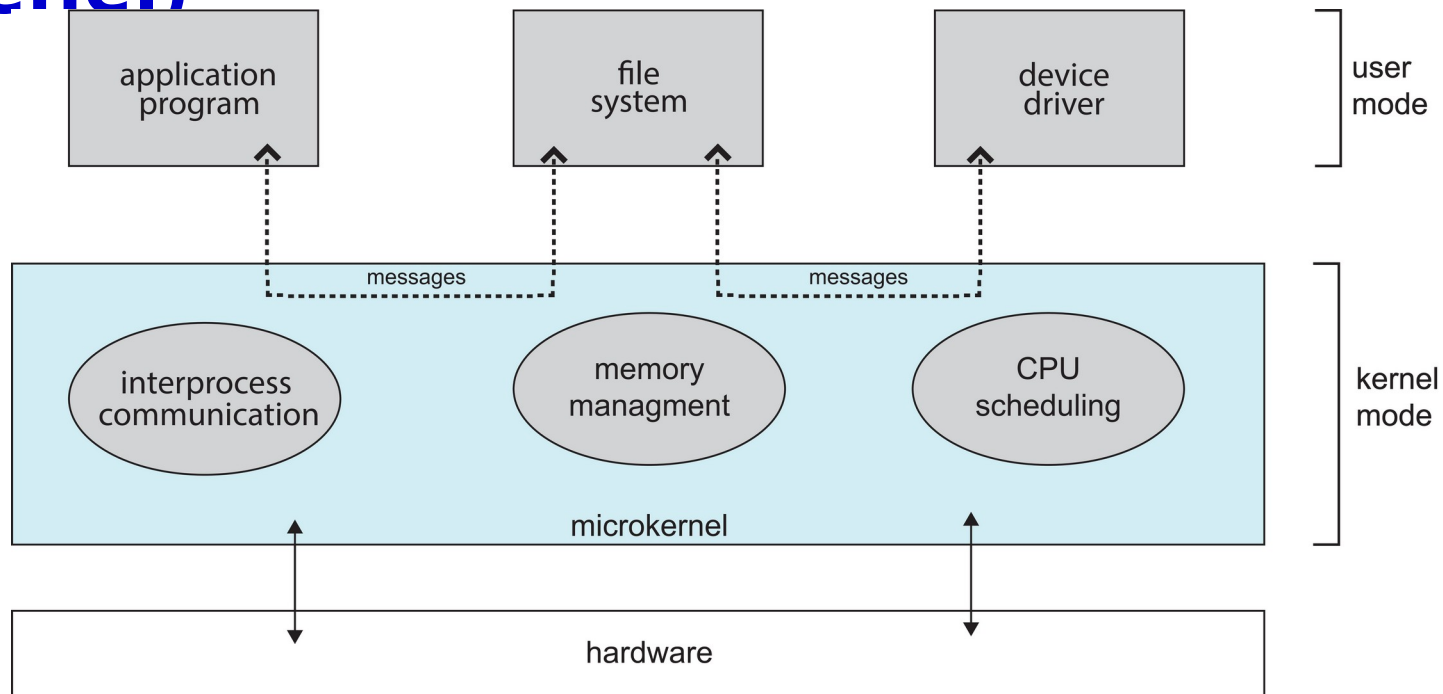
# Layered approach

- **xv6**



# Microkernel approach


- Mach, Mac OS X (based on Mach)
- GNU herd
- Against “monolithic” approach (Linux kernel)



# Skipped

- **IOS, Android structure**

# Clustered sytem



# Free Software / Open Source kernels

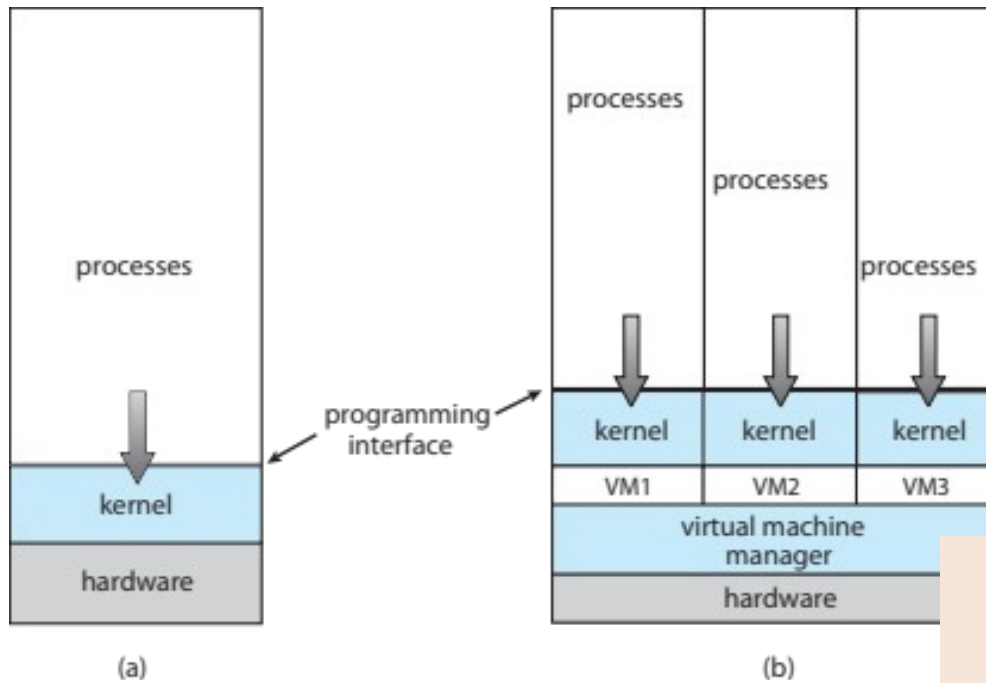
- **Linux**
- **BSD Unix**
- **GNU Herd**
- **xv6**
- **Minix**
- **FreeRTOS**
- **OpenSolaris**
- **Many others**

[https://en.wikipedia.org/wiki/Comparison\\_of\\_open-source\\_operating\\_systems](https://en.wikipedia.org/wiki/Comparison_of_open-source_operating_systems)


# Virtualization

- **A general concept!**
- **Create a “virtual reality” out of “reality”**
  - Black boxes that do something
  - End users need not know “how”, but only “what”
- **Following all are “virtualizations”**
  - Function
  - Class
  - An application
  - Debian GNU/Linux (or any OS)

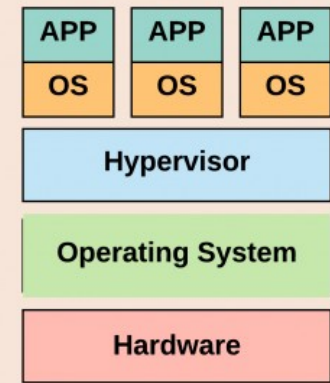
# Virtual Machines, an example



## Types of Hypervisor



Type1 Hypervisor



Type2 Hypervisor