

Inter Process Communication

Revision of process related concepts

- **PCB, struct proc**
- **Process lifecycle – different states**
- **Queues/Lists of processes**
- **What is “Blocking”**
- **Event driven kernel**

Before IPC, let's learn more about file related system calls

- **Redirection**

`ls > /tmp/file`

`cat < /etc/passwd`


how does this work?

- **File descriptors are inherited across fork**

IPC: Inter Process Communication

- Processes within a system may be independent or cooperating
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing, e.g. copy paste
 - Computation speedup, e.g. matrix multiplication
 - Modularity, e.g. chrome – separate process for display, separate for fetching data
 - Convenience
- Cooperating processes need interprocess communication (IPC)
- Broadly Two models of IPC
 - Shared memory (examples: shared memory, pipes, ...)
 - Message passing (examples: send/recv on socket, send-recv messages, ...)

Shared Memory Vs Message Passing



Each requires OS to provide system calls for

- Creating the IPC mechanism
- To read/write using the IPC mechanism
- Delete the IPC mechanism

Note: processes communicating with each other with the help of OS!

Figure 3.11 Communications models. (a) Shared memory. (b) Message passing.

Typical code of shared memory type of solutions

Process P1

```
x = getshm(ID)
```

```
*x = 100;
```

Process P2

```
x = getshm(ID)
```

```
y = *x
```

Typical code of message passing type solutions

Process P1

`send(P2, x)`

or

`broadcast(x)`

Process P2

`y = receive(P1)`

or

`y = receive(anyone)`

Example of co-operating processes: Producer Consumer Problem

- **Paradigm for cooperating processes, producer process produces information that is consumed by a consumer process**
 - **unbounded-buffer places no practical limit on the size of the buffer**
 - **bounded-buffer assumes that there is a fixed buffer size**

Example of co-operating processes: Producer Consumer Problem

- Shared data

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
    . . .
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

- Can only use BUFFER_SIZE-1 elements

Example of co-operating processes: Producer Consumer Problem

▪ Code of Producer

```
while (true) {  
    /* Produce an item */  
    while (((in = (in + 1) % BUFFER SIZE count) ==  
out)  
        ; /* do nothing -- no free buffers */  
    buffer[in] = item;  
    in = (in + 1) % BUFFER SIZE;  
}
```

Example of co-operating processes: Producer Consumer Problem

- **Code of Consumer**

```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```


Pipes

Pipes for IPC

- **Two types**
 - Unnamed Pipes or ordinary pipes
 - Named Pipe

Ordinary pipes

- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the write-end of the pipe)
- Consumer reads from the other end (the read-end of the pipe)
- Ordinary pipes are therefore unidirectional
- Requires a parent-child (or sibling, etc) kind of relationship between communicating processes



Named pipes

- Also called FIFO
- Processes can create a “file” that acts as pipe. Multiple processes can share the file to read/write as a FIFO
- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems
- Is not deleted automatically by OS

Named pipes

- `int mkfifo(const char *pathname, mode_t mode);`
- Example

Shared Memory

System V shared memory

- **Process first creates shared memory segment**
segment id = shmget(IPC_PRIVATE, size, S_IRUSR | S_IWUSR);
- **Process wanting access to that shared memory must attach to it**
shared_memory = (char *) shmat(id, NULL, 0);
- **Now the process could write to the shared memory**
sprintf(shared_memory, "Writing to shared memory");
- **When done, a process can detach the shared memory from its address space**
shmdt(shared_memory);

Example of Shared memory

POSIX Shared Memory

- **What is POSIX?**
 - Portable Operating System Interface (POSIX)
 - family of standards
 - specified by the IEEE Computer Society
 - for maintaining compatibility between operating systems.
 - API (system calls), shells, utility commands for compatibility among UNIXes and variants

POSIX Shared Memory

- **shm_open**
- **ftruncate**
- **Mmap**
- **See the example in Textbook**

Message passing

Message Passing

- **Message system** – processes communicate with each other using `send()`, `receive()` like syscalls given by OS
- **IPC facility provides two operations:**
 - `send(message)` – message size fixed or variable
 - `Receive(message)`
- **If P and Q wish to communicate, they need to:**
 - establish a communication link between them
 - exchange messages via `send/receive`
- **Communication link can be implemented in a variety of ways**

Message Passing using “Naming”

- **Pass a message by “naming” the receiver**
 - **A) Direct communication with receiver**
 - **Receiver is identified by sender directly using it's name**
 - **B) Indirect communication with receiver**
 - **Receiver is identified by sender in-directly using it's 'location of receipt'**

Message passing using direct communication

- **Processes must name each other explicitly:**
 - `send (P, message)` – send a message to process P
 - `receive(Q, message)` – receive a message from process Q
- **Properties of communication link**
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Message passing using IN-direct communication

- **Messages are directed and received from mailboxes (also referred to as ports)**
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- **Properties of communication link**
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Message passing using IN-direct communication

- **Operations**
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox
- **Primitives are defined as:**
 - `send(A, message)` – send a message to mailbox A
 - `receive(A, message)` – receive a message from mailbox A

Message passing using IN-direct communication

- **Mailbox sharing**

- P1, P2, and P3 share mailbox A
- P1, sends; P2 and P3 receive
- Who gets the message?

- **Solutions**

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Message Passing implementation: Synchronization issues

- **Message passing may be either blocking or non-blocking**
- **Blocking is considered synchronous**
 - Blocking send has the sender block until the message is received
 - Blocking receive has the receiver block until a message is available
- **Non-blocking is considered asynchronous**
 - Non-blocking send has the sender send the message and continue
 - Non-blocking receive has the receiver receive a valid message or null

Producer consumer using blocking send and receive

Producer

```
message next produced;  
while (true) {  
    /* produce an item in  
    next_produced */  
    send(next_produced);  
}
```

Consumer

```
message  
next_consumed;  
while (true) {  
    receive(next_consumed);  
}
```

Message Passing implementation: choice of Buffering

- Queue of messages attached to the link; implemented in one of three ways
 1. Zero capacity – 0 messages
 - Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
 - Sender must wait if link full
 3. Unbounded capacity – infinite length
 - Sender never waits

