

# Scheduling

**Abhijit A.M.**  
**[abhijit.comp@coep.ac.in](mailto:abhijit.comp@coep.ac.in)**

**Credits: Slides from [os-book.com](http://os-book.com)**

# Necessity of scheduling

## Multiprogramming

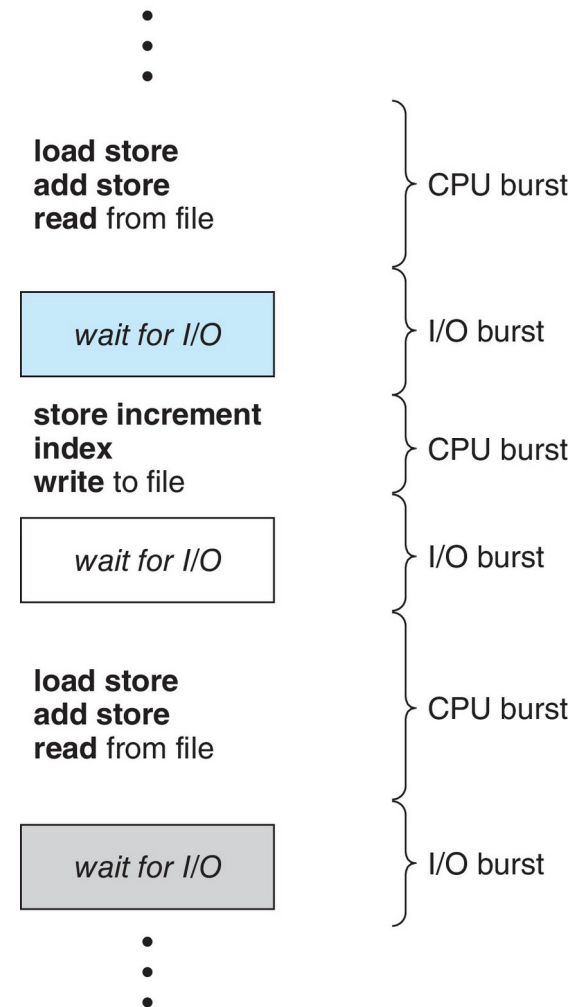
- Ability to run multiple programs at a time
- **Increase use of CPU**
  - CPU utilisation

# CPU Scheduling

- **The task of selecting 'next' process/thread to execute on CPU and doing a context switch**
- **Scheduling algorithm**
  - Criteria for selecting the 'next' process/thread and it's implementation
- **Why is it important?**
  - Affects performance !
  - Affects end user experience !
  - Involves money!

# Observation: CPU, I/O Bursts

- Process can 'wait' for an event (disk I/O, read from keyboard, etc. )
- During this period another process can be scheduled
- **CPU-I/O Burst Cycle:**
  - Process execution consists of a **cycle** of CPU execution and I/O wait
  - **CPU burst** followed by **I/O burst**
  - CPU burst distribution is of main concern



# Let's understand the problem

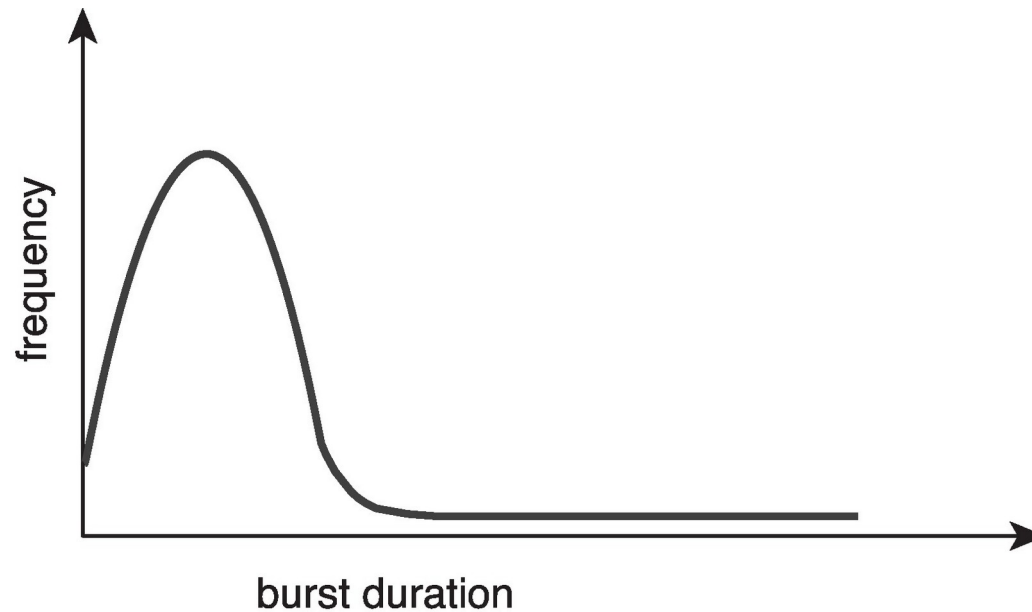
- **Programs have alternate CPU and I/O bursts**

- Some are CPU intensive
- Some are I/O intensive
- Some are mix of both

## A C code example:

```
f(int i, int j, int k) {  
    j = k * i;      // CPU burst  
    scanf("%d", &i); // I/O  
burst  
    k = i * j; // CPU burst  
    printf("%d\n", k); // I/O  
burst  
    return k;  
}
```

# CPU bursts: observation



Large number of short bursts

Small number of longer bursts

# Scheduler, what does it do?

- **From a list of processes, ready to run**
  - **Selects a process for execution**
  - **Allocates a CPU to the process for execution**
  - **Does “context switch”**

# Context and Context Switch

## ▪ Context

- Set of registers. Which ones?
- All or some ?
- Following registers on xv6 kernel: edi, esi, ebx, ebp, eip
  - Related to calling convention!
- Linux kernel context: Much bigger!
- Also includes: MMU setup

## ▪ Context switch

- Process context -> kernel context
  - On interrupt, system call, or exception
- Kernel context -> process context
  - Returning from system call, returning from interrupt handler, scheduling a process
- In every switch, need to change to the set of registers “last in use” by that context, and also the MMU setup




# When is scheduler invoked?

- 1) **Process Switches from running to waiting state**
    - Waiting for I/O, etc.
  - 2) **Switches from running to ready state**
    - E.g. on a timer interrupt
  - 3) **Switches from waiting to ready**
    - I/O wait completed, now ready to run
  - 4) **Terminates**
- **Scheduling under 1 and 4 is nonpreemptive**
  - **All other scheduling is preemptive**
    - Consider access to shared data
    - Consider preemption while in kernel mode
    - Consider interrupts occurring during crucial OS activities

# Dispatcher: A part of scheduler

- Gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- **Dispatch latency**
  - time taken to stop one process and start another running
- **Xv6: `swtch()`, some tail end parts of `sched()`, `trap()`, `trapret()`**



# Dispatcher in action on Linux

- **Run “vmstat 1 3”**
  - Means run vmstat 3 times at 1 second delay
- **In output, look at CPU:cs**
  - Context switches every second
- **Also for a process with pid 3323**
- **Run**  
**“cat /proc/3323/status”**
  - **See**  
voluntary\_ctxt\_switches  
--> Process left CPU  
nonvoluntary\_ctxt\_switches  
--> Process was preempted

# Scheduling criteria

- **CPU utilization: Maximise**
  - keep the CPU as busy as possible. Linux: idle task is scheduled when no process to be scheduled.
- **Throughput : Maximise**
  - # of processes that complete their execution per time unit
- **Turnaround time : Minimise**
  - amount of time to execute a particular process
- **Waiting time : Minimise**
  - amount of time a process has been waiting in the ready queue
- **Response time : Minimise**
  - amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)
-

# To be studied later


- **Evaluation of scheduling criteria**
- **Different scheduling algorithms, and their characteristics**
  - Round Robin, FIFO, Shortest Job First, Priority, Multi-level Queue, etc.
- **Thread scheduling**

# **Multi Processor Scheduling**

# Multiprocessor systems


- **Each processor has separate set of registers**
  - All: eip, esp, cr3, eax, ebx, etc.
- **Each processor runs independently of others**
- **Main difference is in how do they access memory**

# Symmetric multiprocessing (SMP)






# Multicore systems (also SMP)



**No difference from the perspective of OS. The hardware ensures that OS sees multiple processors and not multiple cores.**

# Multicore Processors

- Recent trend to place multiple processor cores on same physical chip
- Faster and consumes less power
- Multiple threads per core also growing
  - Takes advantage of memory stall to make progress on another thread while memory retrieve happens



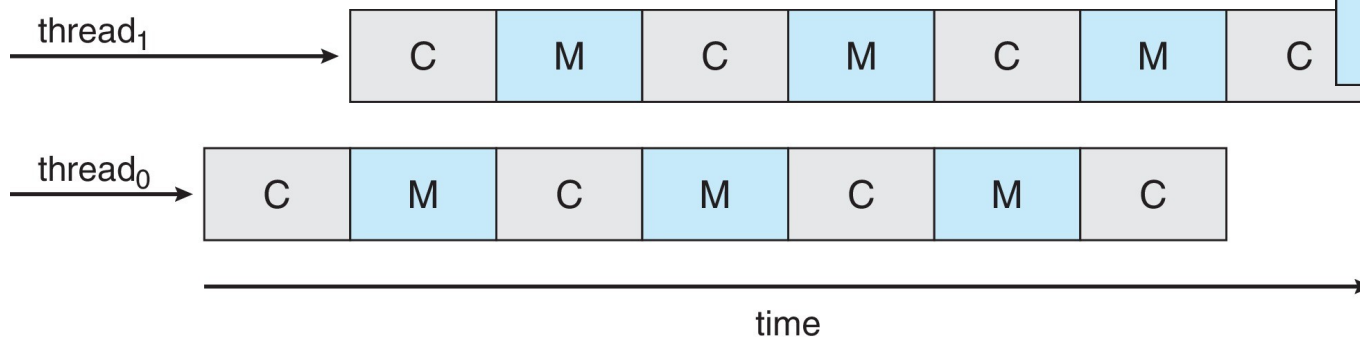

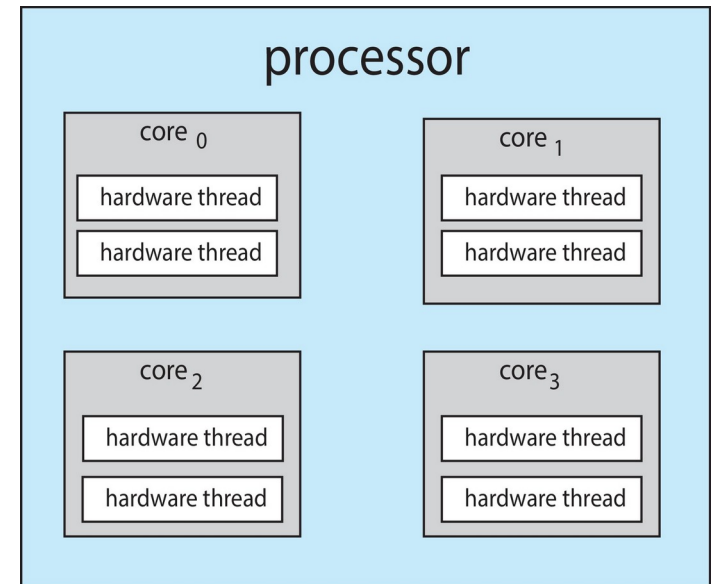
# Multithreaded Multicore System

Each core has  $> 1$  hardware threads.

**Chip-multithreading** (CMT) assigns each core multiple hardware threads. (Intel refers to this as **hyperthreading**.)

On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.

If one thread has a memory stall, switch to another thread!



# More on SMP systems

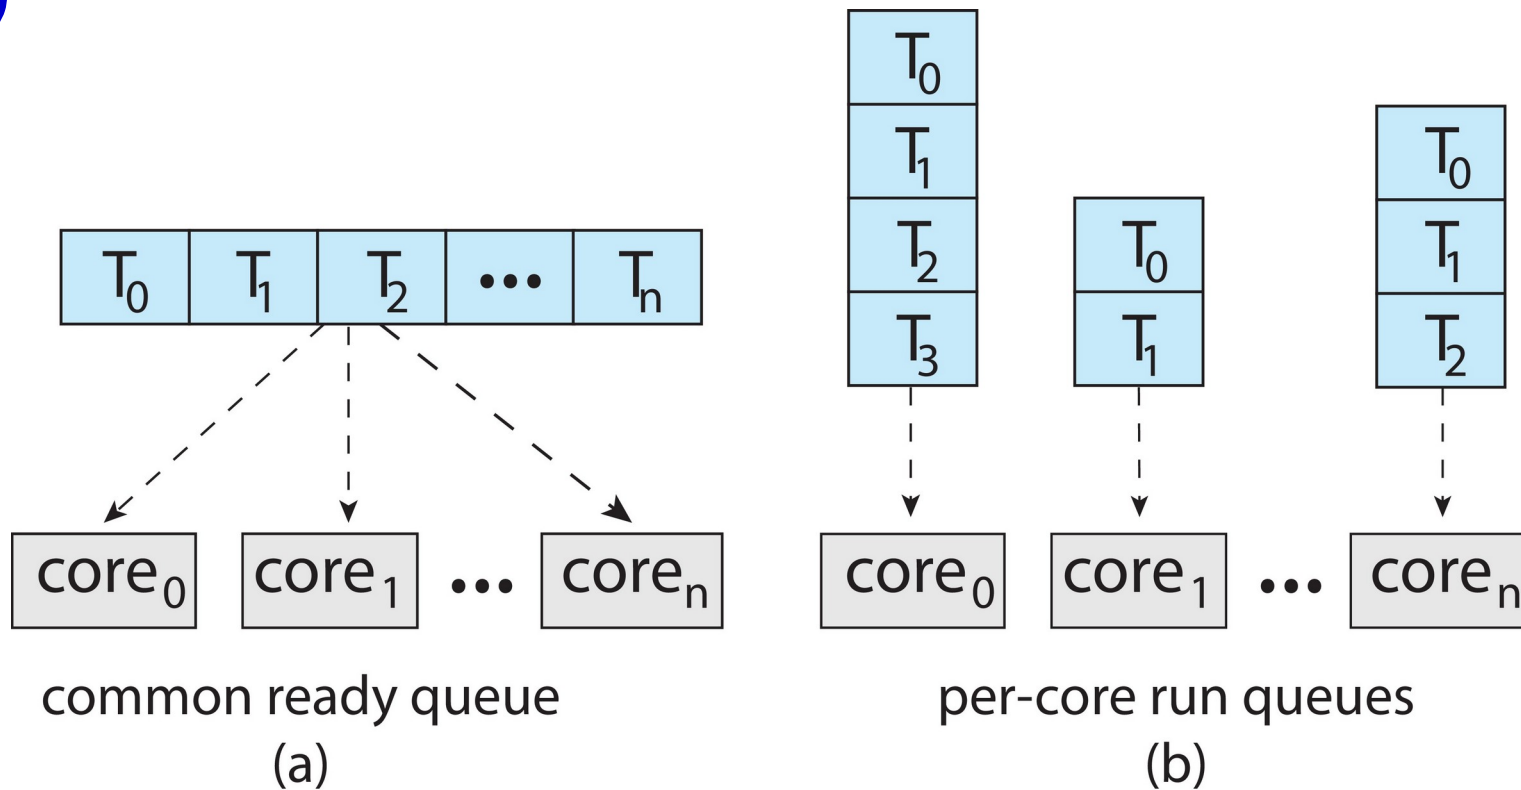
- **During booting – each CPU needs to be turned on**
  - Special I/O instructions writing to particular ports
  - See `lapicstartap()` in xv6
  - Need to setup CR3 on each processor
  - Segmentation, Page tables are shared (same memory for all CPUs)
- **All processors will keep running independently of each other**
- **Different interrupts on each processor – each needs IDT setup**
- **Each processor will be running processes, interrupt handlers, syscalls**
- **Synchronization problems !**
- **How to do scheduling ?**

# Multiple-Processor Scheduling

- **CPU scheduling more complex when multiple CPUs are available**
- **Multiprocess may be any one of the following architectures:**
  - **Multicore CPUs**
  - **Multithreaded cores**
  - **NUMA systems**
  - **Heterogeneous multiprocessing**

# Multiple-Processor Scheduling

- Symmetric multiprocessing (SMP) is where each processor is self scheduling.
- All threads may be in a common ready queue (a)
- Each processor may have its own private queue of threads (b)



# SMP in xv6

- Only one process queue
- No load balancing, no affinity (more later)
- A process may run any CPU-burst /allotted-time-quantum on any processor randomly

**End**