

Threads

Abhijit A M
abhijit.comp@coep.ac.in

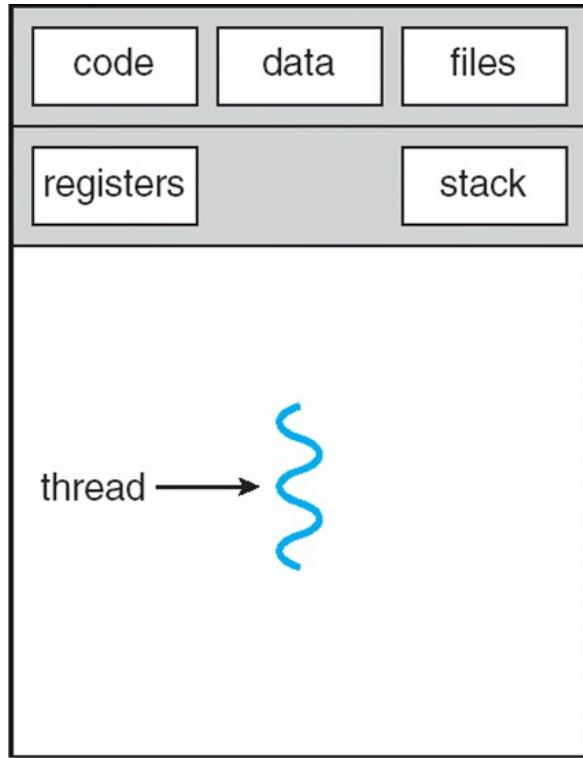
Threads

- **thread — a fundamental unit of CPU utilization**
 - A separate control flow within a program
 - set of instructions that execute “concurrently” with other parts of the code
 - Note the difference: Concurrency: progress at the same time, Parallel: execution at the same time
- **Threads run within application**
 - An application can be divided into multiple parts
 - Each part may be written to execute as a threads
- **Let's see an example**

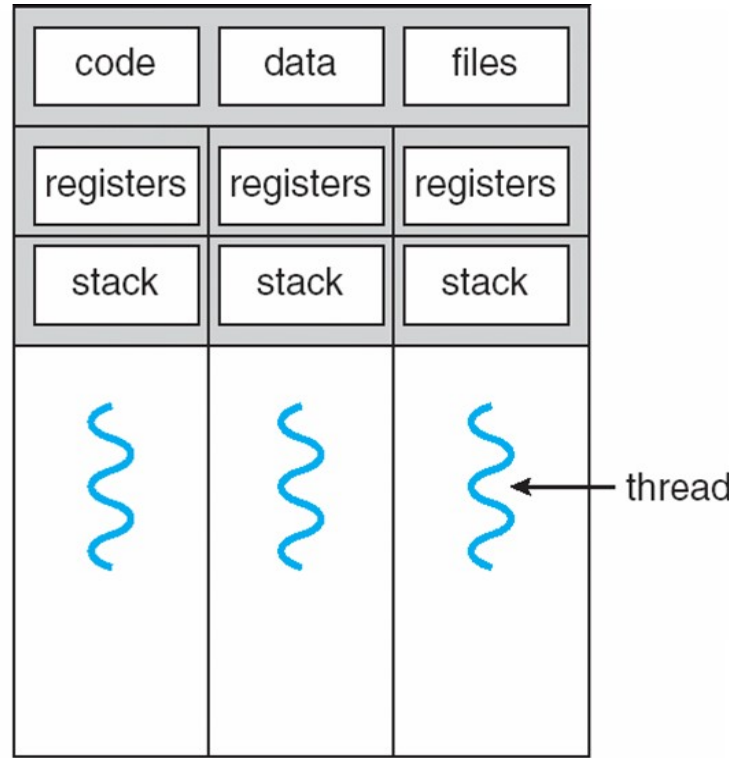
Threads

- **Multiple tasks with the application can be implemented by separate threads**
 - Update display
 - Fetch data
 - Spell checking
 - Answer a network request
- **Process creation is heavy-weight while thread creation is light-weight, due to the very nature of threads**
- **Can simplify code, increase efficiency**
- **Kernels are generally multithreaded**

Single vs Multithreaded process




single-threaded process



multithreaded process


A multithreaded server



Benefits of threads

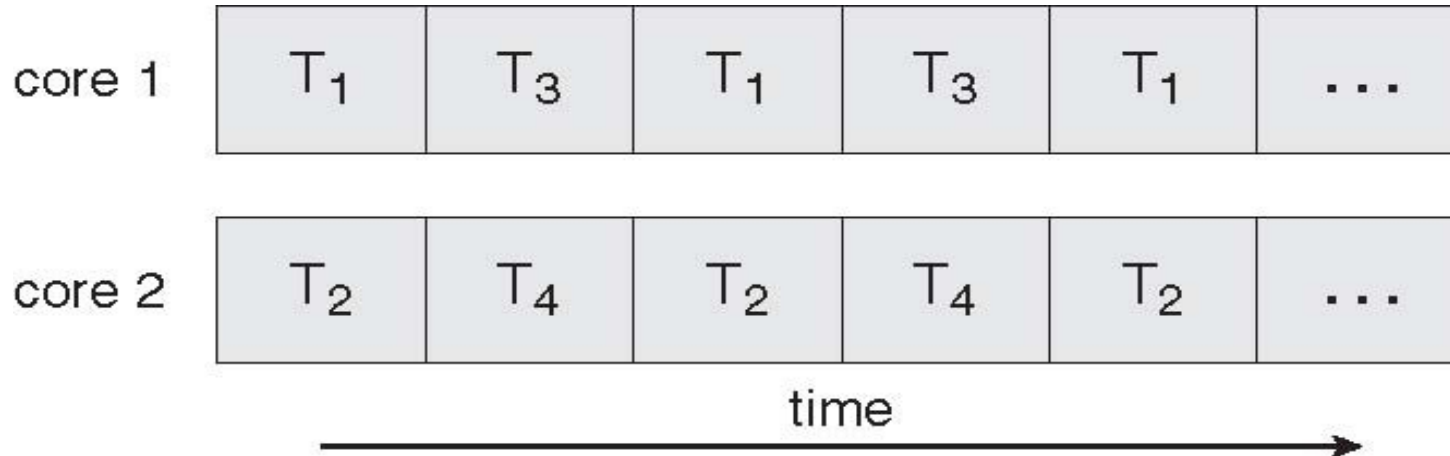
- **Responsiveness**
- **Resource Sharing**
- **Economy**
- **Scalability**

Single vs Multicore systems



Single core : Concurrency possible

Multicore : parallel execution possible



Multicore programming

- **Multicore systems putting pressure on programmers, challenges include:**
 - **Dividing activities**
 - **Balance**
 - **Data splitting**
 - **Data dependency**
 - **Testing and debugging**
 -

User vs Kernel Threads

- **User Threads: Thread management done by user-level threads library**
- **Three primary thread libraries:**
 - POSIX Pthreads
 - Win32 threads
 - Java threads
- **Kernel Threads: Supported by the Kernel**
- **Examples**
 - Windows XP/2000
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X

User threads vs Kernel Threads

▪ User threads

- User level library provides a “typedef” called threads
- The scheduling of threads needs to be implemented in the user level library
- Need some type of timer handling functionality at user level execution of CPU
 - OS needs to provide system calls for this
- Kernel does not know that there are threads!


▪ Kernel Threads

- Kernel implements concept of threads
- Still, there may be a user level library, that maps kernel concept of threads to “user concept” since applications link with user level libraries
- Kernel does scheduling!

Multithreading models

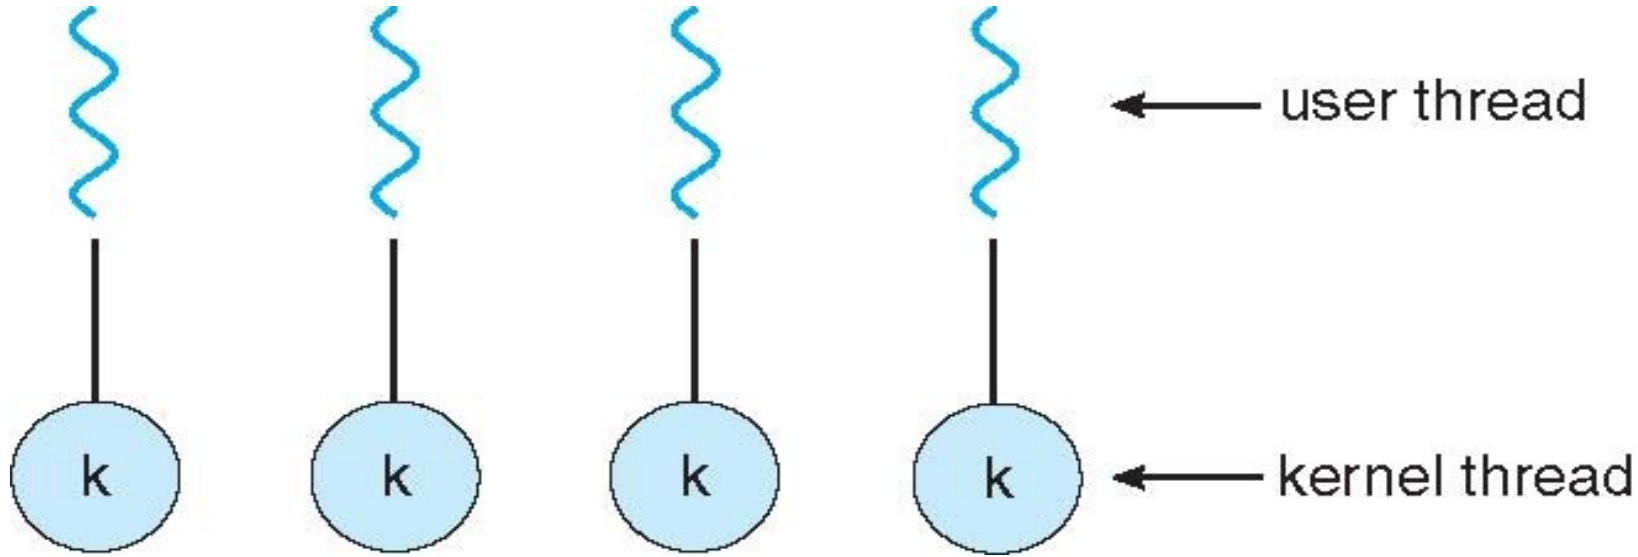
- **How to map user threads to kernel threads?**
 - Many-to-One
 - One-to-One
 - Many-to-Many
- **What if there are no kernel threads?**
 - Then only “one” process. Hence many-one mapping possible, to be done by user level thread library
 - Is One-One possible?

Many-One Model




- **Many user-level threads mapped to single kernel thread**
- **Examples:**
 - **Solaris Green Threads**
 - **GNU Portable Threads**

One-One Model



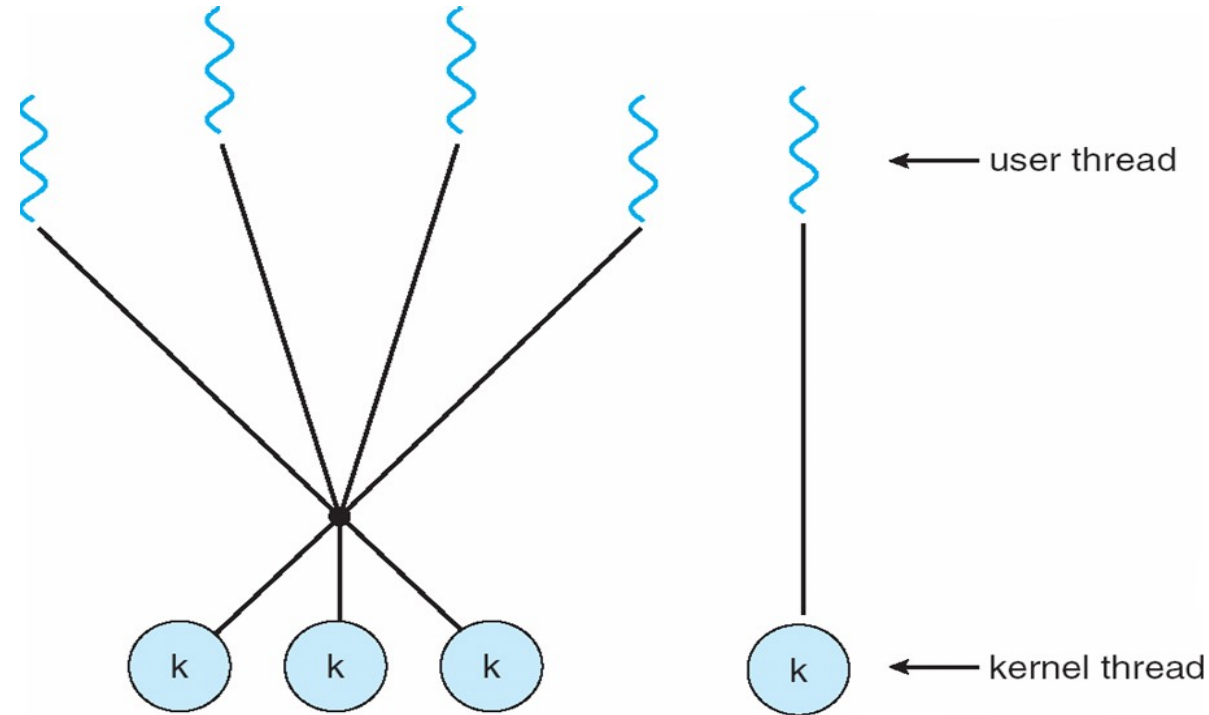
- Each user-level thread maps to kernel thread
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later

Many-Many Model



- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows NT/2000 with the ThreadFiber package

Two Level Model



- Similar to M:M, except that it allows a user thread to be bound to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier

Thread Libraries

Thread libraries

- Thread library provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS

pthread

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)
-

Demo of pthreads code

**Demonstration on Linux – see the code,
compile and execute it.**

Other libraries

- **Windows threading API**
 - **CreateThread(...)**
 - **WaitForSingleObject(...)**
 - **CloseHandle(...)**
- **Java Threads**
 - **The Threads class**
 - **The Runnable Interface**

Issues with threads

- **Semantics of fork() and exec() system calls**
 - Does fork() duplicate only the calling thread or all threads?
- **Thread cancellation of target thread**
 - Terminating a thread before it has finished
 - Two general approaches:
 - Asynchronous cancellation terminates the target thread immediately.
 - Deferred cancellation allows the target thread to periodically check if it should be cancelled.

More on threads

Thread pools

- **Some kernels/libraries can provide system calls to :
Create a number of threads in a pool where they
await work, assign work/function to a waiting thread**
- **Advantages:**
 - **Usually slightly faster to service a request with an
existing thread than create a new thread**
 - **Allows the number of threads in the application(s) to be
bound to the size of the pool**

Thread Local Storage (TLS)

- Thread-specific data, Thread Local Storage (TLS)
 - Not local, but global kind of data for all functions of a thread, more like “static” data
 - Create Facility needed for data private to thread
 - Allows each thread to have its own copy of data
 - Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
 - gcc compiler provides the storage class keyword `thread` for declaring TLS data

```
static __thread int  
threadID;
```

```
int arr[16];  
  
int f() {  
    a(); b(); c();  
}  
  
int g() {  
    x(); y();  
}  
  
int main() {  
    th_create(...,f,...);  
    th_create(...,g,...);  
}  
  
//arr is visible to all of them!  
//need data for only f,a,b,c  
//need data for only g,x,y
```


Thread Local Storage (TLS) in pthreads

- **Functions**

- `pthread_key_create`
- `pthread_key_delete`
- `pthread_setspecific`
- `pthread_getspecific`

- **See**

- `thrd_specific.c` file

Scheduler activations for threads

- **Scheduler Activations**
 - Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
 - Scheduler activations provide upcalls - a communication mechanism from the kernel to the thread library
 - This communication allows an application to maintain the correct number kernel threads

Scheduler activations for threads

Library

```
--  
upcall_handler() {  
    Create one more LWP and schedule threads on this;  
}  
th_setup(int n) {  
    max_LSW = n;  
    curr_LWP = 0;  
    register_upcall(upcall_handler);  
}  
th_create(...., fn,....) {  
    if(curr_LWP < max_LWP)  
        create LWP;  
    schedule fn on one of the LWP;  
}
```


application

```
---  
f() {  
    scanf();  
}  
g() {  
    recv();  
}  
h() {...}; i() {...}  
main()  
    th_setup(2);  
    th_create(....,f,...);  
    th_create(....,g,...);  
    th_create(....,h,...);  
    th_create(....,i,...);  
}
```

Kernel


```
---  
register_upcall(function f) {  
    proc->upcall = f;  
}  
sys_write() {  
    // before calling sleep() going  
    to block  
    myproc()->upcall(); // tricky!  
}
```

Issues with threads




- **Scheduler Activations: LWP approach**
 - An intermediate data structure LWP
 - appears to be a virtual processor on which the application can schedule a user thread to run.
 - Each LWP attached to a kernel thread
 - Typically one LWP per blocking call, e.g. 5 file I/Os in one process, then 5 LWPs needed

Issues with threads



- **Scheduler Activations:
LWP approach**
 - Kernel needs to inform application about events like: a thread is about to block, or wait is over : this is 'upcall'
 - This will help application relinquish the LWP or request a new LWP

Issues with threads



- **Example NetBSD**
 - “An Implementation of Scheduler Activations on the NetBSD Operating System”
 - <https://web.mit.edu/nathanw/www/userix/freenix-sa/freenix-sa.html>

Linux threads

- **Only threads (called task), no processes!**
- **Process is a thread that shares many particular resources with the parent thread**
- **Clone() system call to create a thread**

Linux threads

- **clone()** takes options to determine sharing on process create
- **struct task_struct** points to process data structures (shared or unique depending on clone options)
- **fork()** is a wrapper on top of **clone()**
 - Use 'strace' to see this.

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

Issues in implementing threads project

- How to implement a user land library for threads?
- How to handle 1-1, many-one, many-many implementations?
- Identifying the support required from OS and hardware
- Identifying the libraries that will help in implementation

Issues in implementing threads project

- **Understand the clone() system call completely**
 - Try out various possible ways of calling it
 - Passing different options
 - Passing a user-land buffer as stack
- **How to save and restore context?**
 - C: setjmp, longjmp
 - Setcontext, getcontext(), makecontext(), swapcontext() functions
- **Sigaction is more powerful than signal**
 - Learn SIGALRM handling for timer and scheduler, timer_create() & timer_stop() system calls
- **Customized data structure to store threads, and manage thread-lists for scheduling**

Signals

Signals

- **Signals are used in UNIX systems to notify a process that a particular event has occurred.**
- **Signal handling**
 - Synchronous and asynchronous
- **A signal handler (a function) is used to process signals**
 - Signal is generated by particular event
 - Signal is delivered to a process
 - Then, signal is “handled” by the handler

Signals

- **More about signals**
 - Different signals are typically identified as different numbers
 - Operating systems provide system calls like `kill()` and `signal()` to enable processes to deliver and receive signals
 - `Signal()` - is used by a process to specify a “signal handler” – a code that should run on receiving a signal
 - `Kill()` is used by a process to send another process a signal
 - There are restrictions on which process can send which signal to other processes

Demo

- Let's see a demo of signals with respect to processes
- Let's see signal.h
 - /usr/include/signal.h
 - /usr/include/asm-generic/signal.h
 - /usr/include/linux/signal.h
 - /usr/include/sys/signal.h
 - /usr/include/x86_64-linux-gnu/asm/signal.h
 - /usr/include/x86_64-linux-gnu/sys/signal.h
- man 7 signal
- Important signals: SIGKILL, SIGUSR1, SIGSEGV, SIGALRM, SIGCLD, SIGINT, SIGPIPE, ...

Signal handling by OS

```
Process 12323 {  
    signal(19, abcd);  
}
```

```
OS: sys_signal {  
    Note down that process  
    12323 wants to handle signal  
    number 19 with function  
    abcd  
}
```

```
Process P1 {  
    kill (12323, 19) ;  
}
```

```
OS: sys_kill {
```

Note down in PCB of process 12323 that signal number 19 is pending for you.

```
}
```

When process 12323 is scheduled, at that time the OS will check for pending signals, and invoke the appropriate signal handler for a pending signal.

Threads and Signals

- **Signal handling Options:**
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process