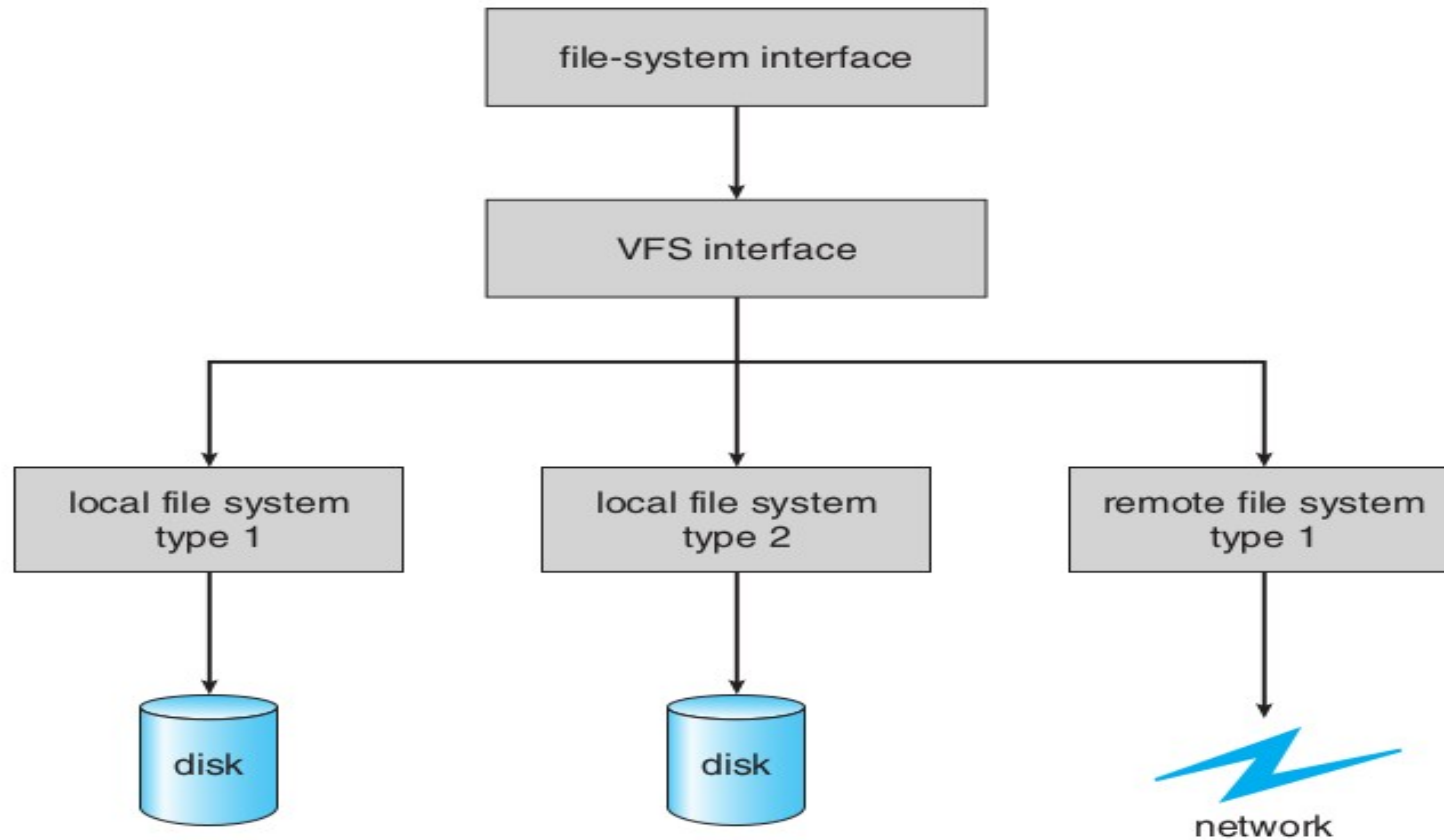


**VFS**



**Figure 15.5** Schematic view of a virtual file system.

# VFS

- **Consider this**

`/dev/sda1` is “/”

`/dev/sda2` is mounted on “/a/b” folder

How does this work in kernel?

`open(“/a/b/c/d”, O_RDONLY)`

- **Consider xv6 code**

- `sys_open` -> `namei` -> `namex` -> (`skipelem`, `dirlookup`, `ilock`)
- `Dirlookup()` of “c” in “/a/b” should return : Not the inode of “c” on `/dev/sda1` but inode of “/” on `/dev/sda2`

# VFS

- **Object Oriented Programming in C (let's see example of this)**
  - Clever use of function pointers
- **There is an “abstract” file system class (VFS), and there are concrete file system classes (ext2, vfat, ...)**
  - `sys_read` → `fileread` → `readi()` becomes
  - `sys_read` → `fileread` → `(i->i_ops->read)()`
- **Inode is a generic inode**
  - Contains file system specific inode pointer
  - And file system specific inode operations
  - Fields setup during `namei()`

```
struct inode_operations {  
    int (*readi) (int, char *, int);  
    int (*writei) (int, char *, int);  
    ....  
}  
struct inode {  
    int mode,  
    int uid;  
    ....  
    void *inode_specific;  
    struct inode_ops i_ops;  
}
```

**Efficiency and Performance  
(and the risks created  
while trying to achieve it!)**

# Efficiency

- **Efficiency dependent on:**
  - Disk allocation and directory algorithms
  - Types of data kept in file's directory entry
  - Pre-allocation or as-needed allocation of metadata structures
  - Fixed-size or varying-size data structures
  -

# Performance

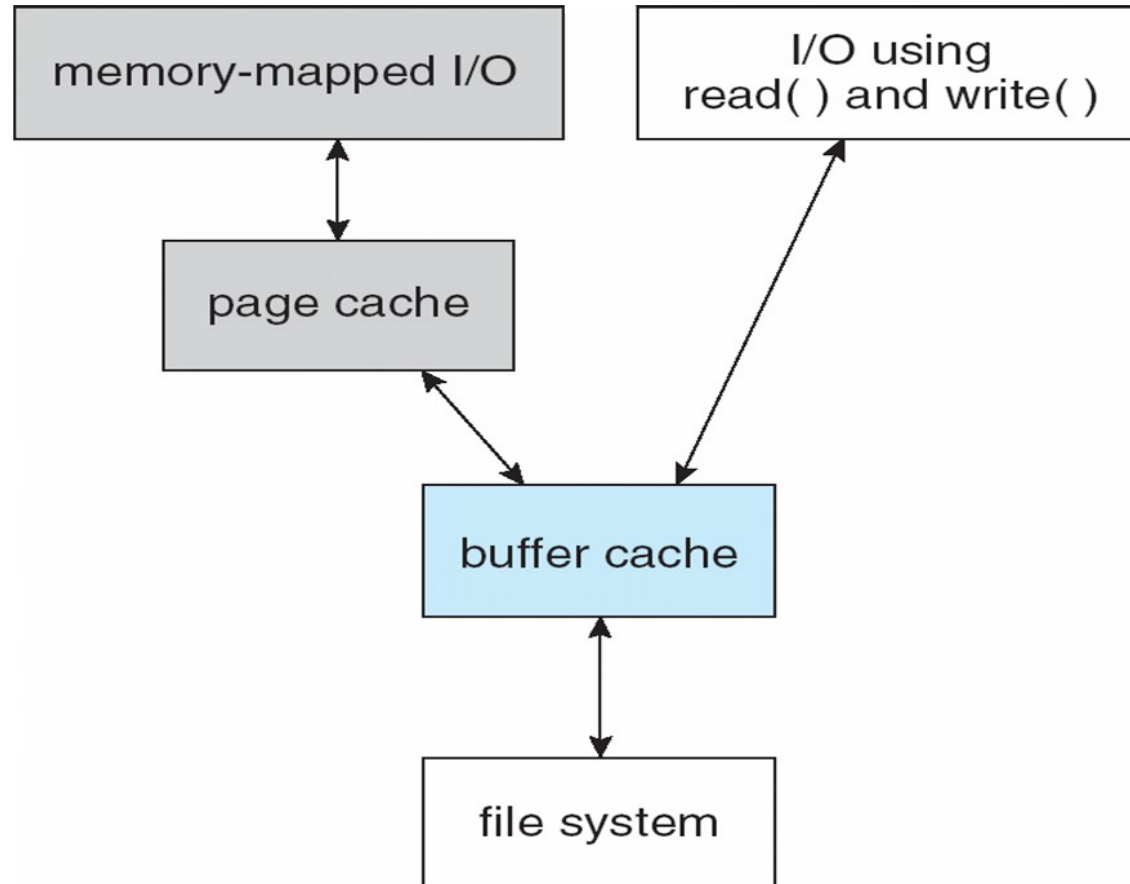
- Keeping data and metadata close together
- Buffer cache – separate section of main memory for frequently used blocks
- Synchronous writes sometimes requested by apps or needed by OS
- No buffering / caching – writes must hit disk before acknowledgement
- Asynchronous writes more common, buffer-able, faster
- Free-behind and read-ahead – techniques to optimize sequential access
- Reads frequently slower than writes

# Page cache

- A page cache caches pages rather than disk blocks using virtual memory techniques and addresses
- Memory-mapped I/O uses a page cache
- Routine I/O through the file system uses the buffer (disk) cache
- This leads to the following figure



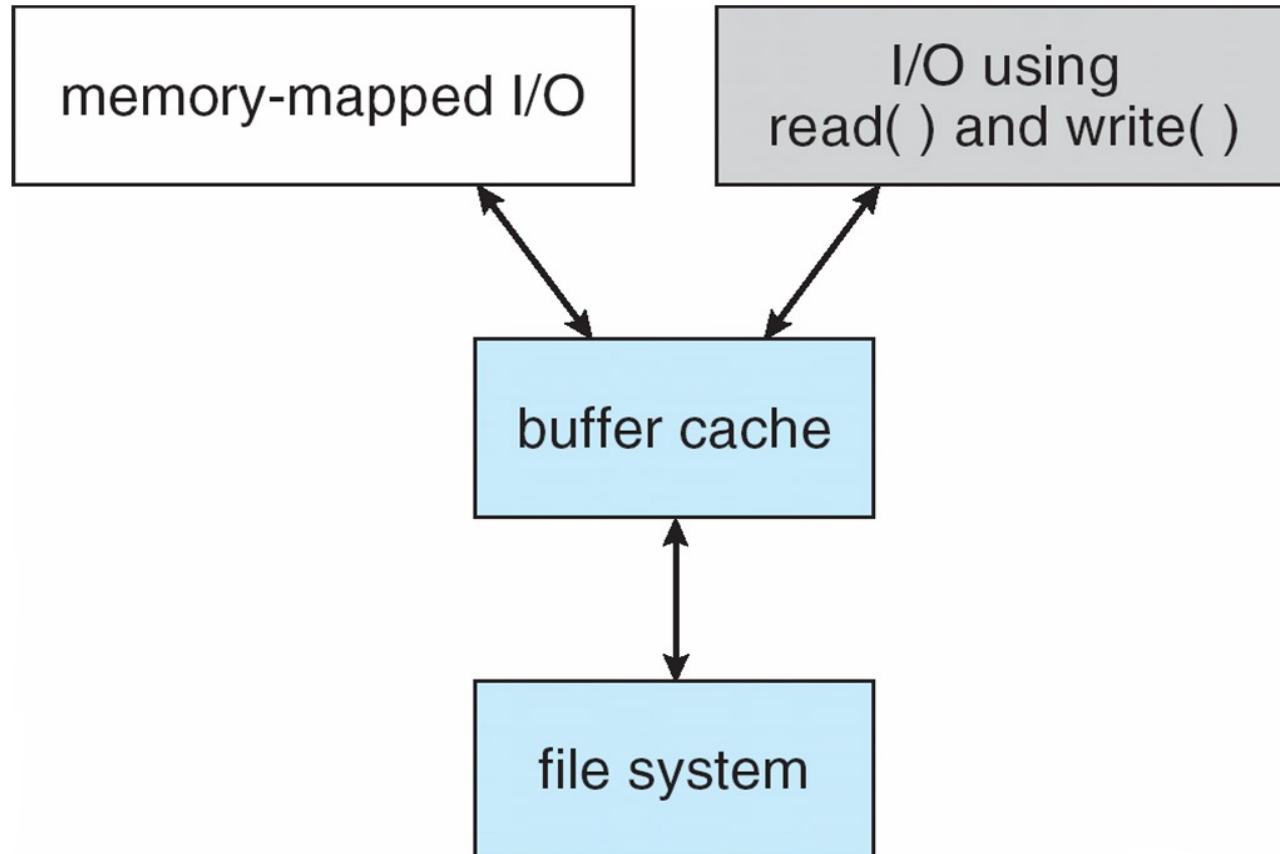
# I/O Without a Unified Buffer Cache



# Unified buffer cache

- A unified buffer cache uses the same page cache to cache both memory-mapped pages and ordinary file system I/O to avoid double caching
- But which caches get priority, and what replacement algorithms to use?

# I/O Using a Unified Buffer Cache



# Recovery

- **Problem. Consider creating a file on ext2 file system.**
  - Following on disk data structures will/may get modified
  - Directory data block, new directory data block, block bitmap, inode table, inode table bitmap, group descriptor, super block, data blocks for new file, more data block bitmaps, ...
  - All cached in memory by OS
- **Delayed write – OS writes changes in its in-memory data structures, and schedules writes to disk when convenient**
  - Possible that some of the above changes are written, but some are not
  - Inconsistent data structure! --> Example: inode table written, inode bitmap written, but directory data block not written

# Recovery

- **fsck: Consistency checking – compares data in directory structure with data blocks on disk, and tries to fix inconsistencies**
  - **Can be slow and sometimes fails**
- **Use system programs to back up data from disk to another storage device (magnetic tape, other magnetic disk, optical)**
- **Recover lost file or disk by restoring data from backup**
- **Faster recovery? - “log structured file system” or “journaling file system” can help**

# Log structured file systems

- Log structured (or journaling) file systems record each metadata update to the file system as a transaction
- All transactions are written to a log
  - A transaction is considered committed once it is written to the log (sequentially)
  - Sometimes to a separate device or section of disk
  - However, the file system may not yet be updated
- The transactions in the log are asynchronously written to the file system structures
  - When the file system structures are modified, the transaction is removed from the log
- If the file system crashes, all remaining transactions in the log must still be performed
- Faster recovery from crash, removes chance of inconsistency of metadata

# Journaling file systems

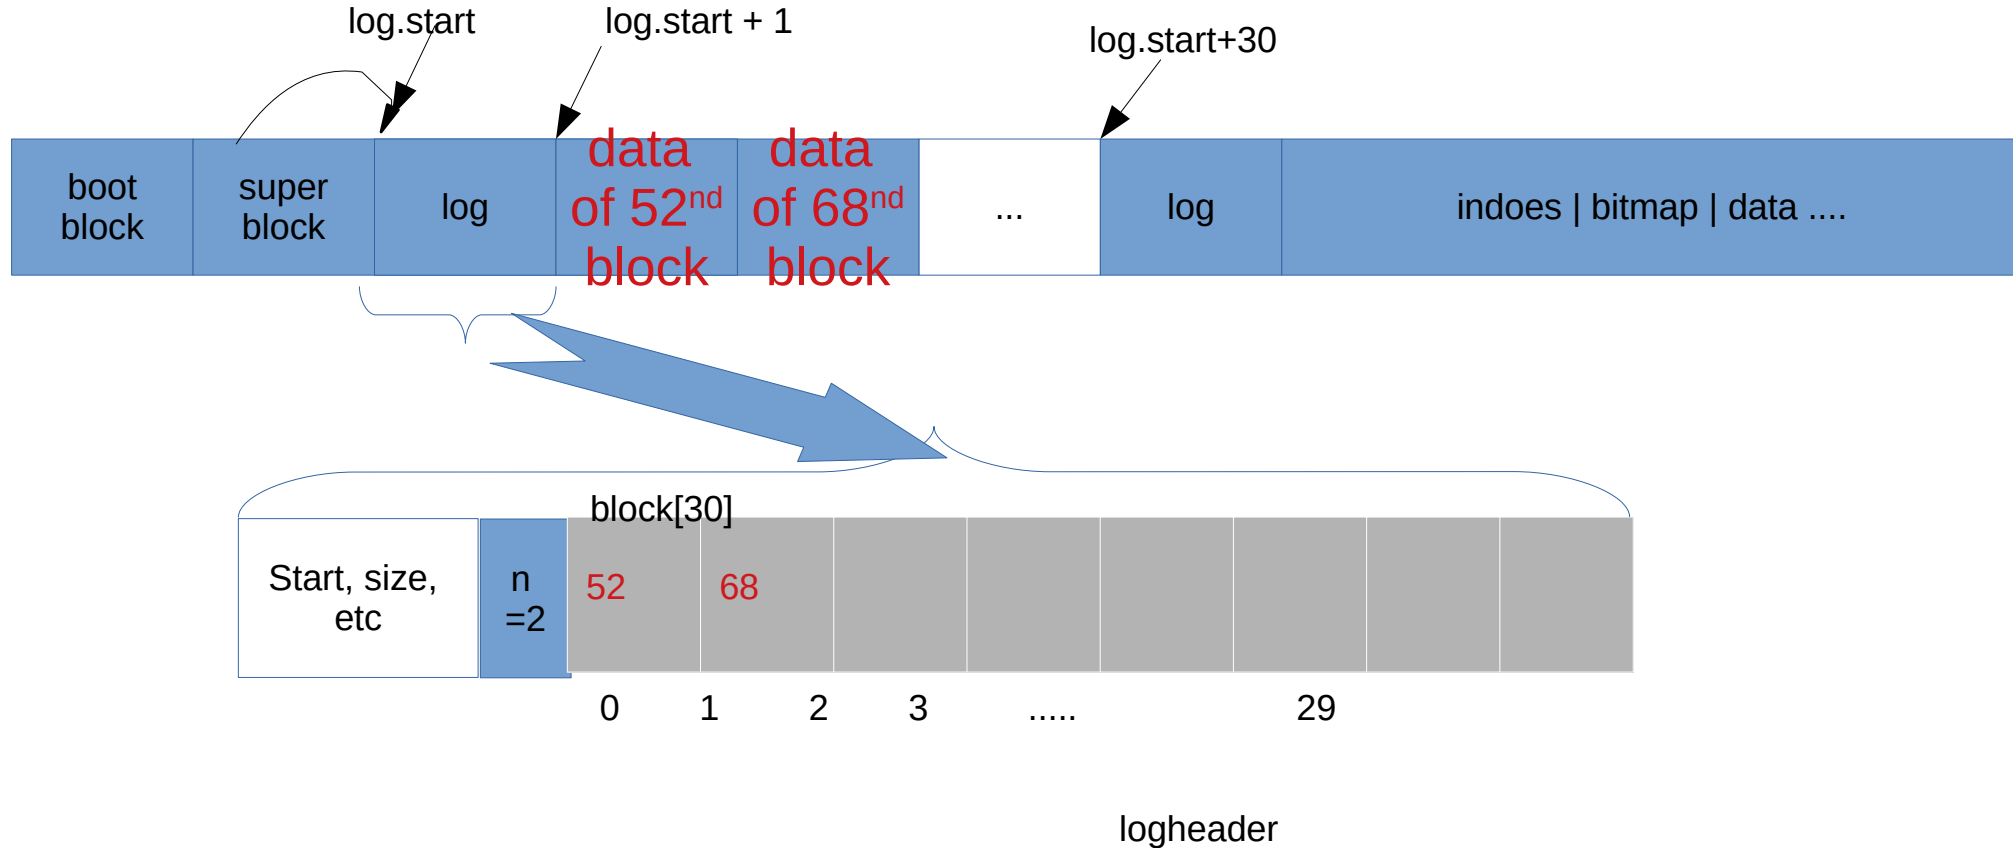
- Veritas FS
- Ext3, Ext4
- Xv6 file system!

# log in xv6

- a mechanism of recovery from disk
- **Concept: multiple write operations needed for system calls (e.g. 'open' system call to create a file in a directory)**
  - some writes succeed and some don't
  - leading to inconsistencies on disk
- **In the log, all changes for a 'transaction' (an operation) are either written completely or not at all**
- **During recovery, completed operations can be "rerun" and incomplete operations neglected**



# log on disk



# log in xv6

- **xv6 system call does not directly write the on-disk file system data structures.**
- **A system call calls `begin_op()` at beginning and `end_op()` at end**
  - `begin_op()` increments `log.outstanding`
  - `end_op()` decrements `log.outstanding`, and if it's 0, then calls `commit()`
- **During the code of system call, whenever a buffer is modified, (and done with)**
  - `log_write()` is called
  - This copies the block in an array of blocks inside log, the block is not written in its actual place in FS as of now
- **when finally `commit()` is called, all modified blocks are copied to disk in the file system**

# log

```
struct logheader { // ON DISK
    int n; // number of entries in use in block[] below
    int block[LOGSIZE]; // List of block numbers stored
};

struct log { // only in memory
    struct spinlock lock;
    int start; // first log block on disk (starts with logheader)
    int size; // total number of log blocks (in use out of 30)
    int outstanding; // how many FS sys calls are executing.
    int committing; // in commit(), please wait.
    int dev; // FS device
    struct logheader lh; // copy of the on disk logheader
};

struct log log;
```

# Typical use case of logging

```
/* In a system call code */  
begin_op();  
...  
bp = bread(...);  
bp->data[...] = ...;  
log_write(bp);  
...  
end_op();
```

prepare for logging. Wait if logging system is not ready or 'committing'. ++outstanding

read and get access to a data block – as a buffer

modify buffer

note down this buffer for writing, in log. proxy for bwrite(). Mark B\_DIRTY. Absorb multiple writes into one.

Syscall done. write log and all blocks. --outstanding.

If outstanding = 0, commit().

# Example of calls to logging

```
//file_write() code
begin_op();
ilock(f->ip);
    /*loop */ r = writei(f-
>ip, ...);
iunlock(f->ip);
end_op();
```

- each writei() in turn calls bread(), log\_write() and brelse()
  - also calls iupdate(ip) which also calls bread, log\_write and brelse
- Multiple writes are combined between begin\_op() and end\_op()

# Logging functions

- **Initlog()**
  - Set fields in global **log.xyz** variables, using FS superblock
  - Recovery if needed
  - **Called from first forkret()**
- **Following three called by FS code**
- **begin\_op(void)**
  - Increment **log.outstanding**
- **end\_op(void)**
  - Decrement **log.outstanding** and call **commit()** if it's zero
- **log\_write(buf \*)**
  - Remember the specified block number in **log.lh.block[]** array
  - Set the block to be dirty
- **write\_log(void)**
  - **Called only from commit()**
  - Use block numbers specified in **log.lh.block** and copy those blocks from memory to log-blocks
- **commit(void)**
  - **Called only from end\_op()**
  - **write\_log()**
  - Write header to disk log-header
  - Copy from log blocks to actual FS blocks
  - Reset and write log header again













