

Scheduling Algorithms

Abhijit A.M.
abhijit.comp@coep.ac.in

Credits: Slides from os-book.com

Calculations of different scheduling criteria

- **If you want to evaluate an algorithm practically, you need a proper workload !**
 - Processes with CPU and I/O bursts
 - Different durations of CPU bursts
 - Different durations of I/O bursts
 - How to do this programmatically?
 - How to ensure that after 2 seconds an I/O takes place?
 - Need periods when system will be “idle” – no process schedulable !

Calculations of different criteria

- **CPU Utilization**

- % time spent in doing 'useful' work
- What is useful work?
 - **On linux**
 - there is an "idle" thread, scheduled when no other task is RUNNABLE
 - Not running idle thread is productive work
 - Includes process + scheduling time + interrupts
 - **On other systems?**
 - Need to define
 - **On xv6**
 - We can say that time spent in the loop selecting a process is idle work

Calculations of different criteria

- **Throughput**

- # processes that complete execution per unit time
- Formula: total # processes completed / total time
- Simply divide by your total workload that completed by the time taken
- Depends on the workload as well. 'long' or 'short' processes.
- If too many short processes , then throughput may appear to be high, like 10s of processes per second

Calculations of different criteria

- **Turnaround time**

- Amount of time required for one process to complete
- For every process, note down the starting and ending time, difference is TA-time
- For process P1 : (Time when process ended – time when process started)
= Sum of time spent in (ready queue + running + waiting for I/O)
- One can find the average T.A. time

Calculations of different criteria

- **Waiting time**
 - amount of time a process has been waiting in the *ready* queue.
 - **To be minimised.**
 - **Part of Turn Around time**
 - **CPU scheduling does not affect waiting time in I/O queues, it affects time in ready queue**

Scheduling Criteria

- **Response time**
 - amount of time it takes from when a request was submitted until the first response (not full output) is produced, (for time-sharing environment).
 - **To be minimised.**
 - E.g. time between your press of a key , and that key being shown in screen

Challenges in implementing the scheduling algorithms

- **Not possible to know number of CPU and I/O bursts and the duration of each before the process runs !**
 - **Although when we do numerical problems around each algorithm, we assume some values for CPU and I/O bursts, so the problems are solved in “hindsight” !**
-

GANTT chart

- A timeline chart showing the sequence in which processes get scheduled
- Used for analysing a scheduling algorithm




Scheduling Algorithms

First- Come, First-Served (FCFS) Scheduling

Burst Time	Process
	P1
	24
3	P2
3	P3

Suppose that the processes arrive in the order: P1 , P2 , P3

The Gantt Chart for the schedule is:



Waiting time for P1 = 0; P2 = 24; P3 = 27

Average waiting time: $(0 + 24 + 27)/3 = 17$

Non Pre-emptive algorithm

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

P2 , P3 , P1

The Gantt chart for the schedule is:



Waiting time for P1 = 6; P2 = 0; P3 = 3

Average waiting time: $(6 + 0 + 3)/3 = 3$

Much better than previous case

FCFS: Convoy effect

- Consider one CPU-bound and many I/O-bound processes
- CPU bound process is scheduled, I/O bound processes are waiting in I/O queues
- I/O bound processes finish I/O and move to ready queue, and wait for CPU bound process to finish
 - I/O devices Idle
- CPU bound process over, goes for I/O. I/O bound processes run quickly, move to I/O queues again
 - CPU idle
- CPU bound process will run when it's ready to run
- Same process will repeat
- --> Lower CPU utilisation
 - Better if I/O bound processes run first


FCFS: further evaluation

- **Troublesome for interactive processes**
 - CPU bound process may hog CPU
 - Interactive process may not get a chance to run early and response time may be quite bad

Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time. Better name – **Shortest Next CPU Burst Scheduler**
- SJF is optimal – gives minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request
 - Could ask the user – bad idea, unlikely to know!

Example of SJF



$$\text{Average waiting time} = (3 + 16 + 9 + 0) / 4 = 7$$

Determining Length of Next CPU Burst


- Not possible to implement SJF as can't know “next” CPU burst. Can only estimate the length – should be similar to the previous one
 - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, using exponential averaging
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1$
 4. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.
- Commonly, α set to $\frac{1}{2}$
- Preemptive version called shortest-remaining-time-first

Examples of Exponential Averaging

- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count
- $\alpha = 1$
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts
- If we expand the formula, we get:
 - $$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots$$
$$+ (1 - \alpha)^j \alpha t_{n-j} + \dots$$
$$+ (1 - \alpha)^{n+1} \tau_0$$
- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Prediction of the Length of the Next CPU Burst

- $\alpha = 1/2$, $\tau_0 = 10$

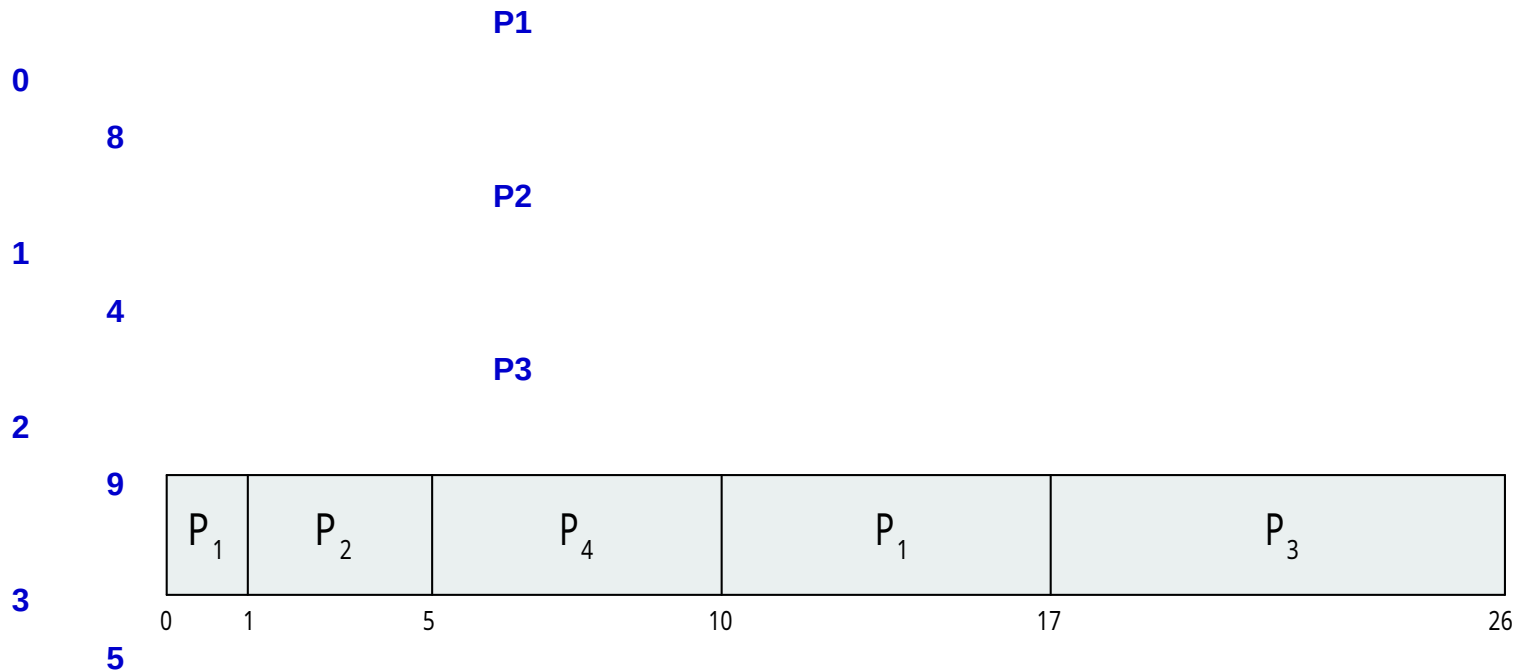


CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

Example of Shortest-remaining-time-first

Preemptive SJF = SRTF. Now we add the concepts of varying arrival times and preemption to the analysis

Process
Arrival Time
Burst Time



Preemptive SJF Gantt Chart

Average waiting time = $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5$ msec

Round Robin (RR) Scheduling

- **Each process gets a small unit of CPU time (time quantum q), usually 10-100 milliseconds.**
 - After this time has elapsed, the process is preempted and added to the end of the ready queue.
- **If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once.**
 - No process waits more than $(n-1)q$ time units.

Round Robin (RR) Scheduling

- **Timer interrupts every quantum to schedule next process**
- **Performance**
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 4

Process
Burst Time

P_1

24


P_2

3

P_3

3

The Gantt chart is:




Typically, higher average turnaround than SJF, but better *response*


q should be large compared to context switch time

q usually 10ms to 100ms, context switch < 10 usec

Time Quantum and Context Switch Time



Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

80% of CPU
bursts should
be shorter
than quantum


Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive (timer interrupt, more time for more priority)
 - Nonpreemptive (no timer interrupt, just schedule process with highest priority)
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Problem \equiv Starvation – low priority processes may never execute
- Solution \equiv Aging – as time progresses increase the priority of the process
-

Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1 10		3
P_2 4	1	1
P_3 4		2
P_4 2	5	1
P_5 2		5

- Priority scheduling Gantt Chart




- Average waiting time = 8.2 msec

Priority Scheduling with Round-Robin

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	4	3
P_2	5	2
P_3	8	2
P_4	7	1
P_5	3	3



Run the process with the highest priority. Processes with the same priority run round-robin

Gantt Chart with 2 ms time quantum




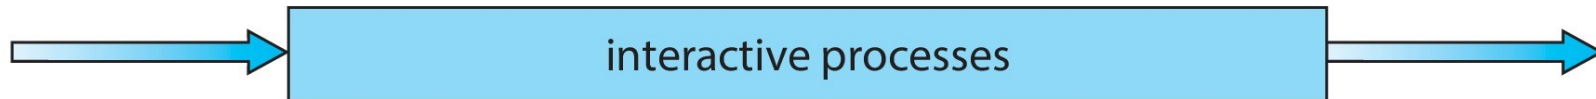
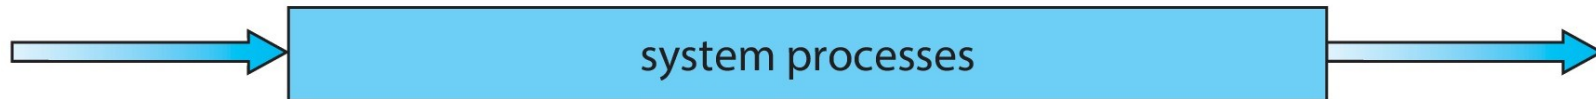
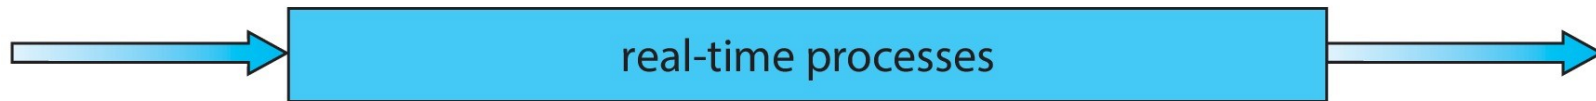
Multilevel Queue

- With priority scheduling, have separate queues for each priority.
- Schedule the process in the highest-priority queue!



Multilevel Queue

highest priority



lowest priority

Implementing multilevel queue


- **Processes need to have a priority**
 - Either modify `fork()/exec()` to have a priority
 - Or add a `nice()` system call to set priority
- **How to know the priority?**
 - The end user of the computer system needs to know this from needs of real life
 - E.g. on a database system, the database process will have a higher priority than other processes

Multilevel Feedback Queue

- **A process can move between the various queues; aging can be implemented this way**
- **Multilevel-feedback-queue scheduler defined by the following parameters:**
 - **number of queues**
 - **scheduling algorithms for each queue**
 - **method used to determine when to upgrade a process**
 - **method used to determine when to demote a process**
 - **method used to determine which queue a process will enter when that process needs service**

Example of Multilevel Feedback Queue


- **Three queues:**
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- **Scheduling rules**
 - Serve all processes in Q_0 first
 - Only when Q_0 is empty, serve processes in Q_1
 - Only when Q_0 and Q_1 are empty, serve processes in Q_2



Example of Multilevel Feedback Queue

- **Scheduling**

- A new job enters queue Q_0
 - When it gains CPU, job receives 8 milliseconds
 - If it does not finish in 8 milliseconds, job is moved to queue Q_1
- At Q_1 job receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2
- To prevent starvation, move a process from lower-priority queue to higher priority queue after it has waited for too long



Thread Scheduling

- Distinction between user-level and kernel-level threads
- When threads supported, threads scheduled, not processes
- Many-to-one and many-to-many models, thread library schedules user-level threads to run on LWP
 - Known as **process-contention scope (PCS)** since scheduling competition is within the process
 - Typically done via priority set by programmer
- Kernel thread scheduled onto available CPU is **system-contention scope (SCS)** – competition among all threads in system

Pthread Scheduling

- **PTHREAD_SCOPE_PROCESS** schedules threads using PCS scheduling
- **PTHREAD_SCOPE_SYSTEM** schedules threads using SCS scheduling
- Linux and macOS only allow **PTHREAD_SCOPE_SYSTEM**
- Let's see a Demo using a program

Multiple-Processor Scheduling – Load Balancing

- **Soft affinity** – the operating system attempts to keep a thread running on the same processor, but no guarantees.
- **Hard affinity** – allows a process to specify a set of processors it may run on.

End